

When and Why Distribution-Free and Generative Quantile Models

Beat the Industry Standard for Downside Risk*

Sean Qin

July 2026 — Draft v2

Abstract

Parametric volatility models with a fixed innovation law remain the backbone of industry downside-risk measurement, while banks' regulatory practice under the Basel Fundamental Review of the Trading Book (FRTB) rests on historical and filtered historical simulation evaluated at the 97.5% Expected Shortfall (the average loss beyond the Value-at-Risk threshold). We ask when, why, and by how much distribution-free and generative conditional quantile models beat this industry standard. Our organizing empirical result is a *misspecification frontier*: a live, per-asset, per-day score — the excess kurtosis (tail thickness beyond the assumed law) and asymmetry of recent GARCH-standardized residuals — that predicts, with forecast-comparison (Diebold–Mariano) significance date by date, when nonparametric conditional quantiles beat GARCH- t . In the top decile of the score the nonparametric edge is +2.71% of pinball loss (DM 6.5); over the remaining ninety percent of asset-days the parametric model is as good or better. The winning nonparametric forecaster is a single *amortized* model trained across hundreds of names — one fit replaces per-asset estimation, beats own-history benchmarks at every listing age, and prices assets with no return history at all, the capability per-asset GARCH cannot offer. The same axis organizes the geography of the edge across four universes: US equities, FX (hyperinflation currencies are the only decisive standalone wins, 12–14%, tier-pooled DM 4.1), twenty-six country equity indices (a developed→frontier gradient with correlation 0.53 to residual kurtosis), and forty-three cross-asset instruments (correlation 0.43). Exploiting the frontier, a residual-hybrid engine — GARCH scale $\times$ state-conditioned nonparametric residual quantiles, finished with an extreme-value-theory tail and a split-conformal (distribution-free, finite-sample) recalibration — survives the Model Confidence Set over the full regulatory battery jointly with only one rival (CAViaR), beats GARCH- t , filtered historical simulation, exponential smoothing, historical simulation and GJR-skew- t with Diebold–Mariano significance, delivers the best-calibrated 97.5% Expected Shortfall, and is the only entrant passing the regulatory exception-count backtests at both levels; the edge grows at the ten-day regulatory horizon,

*Working paper, first full draft July 2026. This draft supersedes the July 12 GRAFT-Q draft, which is retained as the source for the trading application appendix. All results are out-of-sample under walk-forward protocols described in the text; result files and reproducible code are archived in the project repository. A live forward holdout began July 2026 and will be reported regardless of outcome.

where GARCH $\sqrt{h}$ -scaling over-states Expected Shortfall and hence over-charges capital. The generative members of the family then deliver two things no point-forecasting model provides. First, calibrated uncertainty bands on the risk number itself: block-bootstrap and extreme-value-theory intervals on GARCH residuals attain 0.89–0.95 coverage against a 0.90 target for (VaR, ES) — and we show why residual space is essential, measuring that a naive loss-based (Gibbs) posterior on raw returns understates VaR uncertainty by a factor of 1.6. Second, Bayesian posteriors for volatility models whose likelihoods are intractable — fitted purely by simulating the model and validated by simulation-based calibration — including recovery of the roughness (Hurst) parameter of rough volatility (42% prior-uncertainty reduction, calibration coverage 0.91). Honest negatives are reported throughout: multivariate GARCH with dynamic conditional correlations wins equity portfolio tails against three direct nonparametric attempts (a covariance-scale/nonparametric-shape hybrid rescues the deep 1% co-crash tail); a learned regime gate adds nothing over always using the amortized model; and gradient-boosted trees, not the neural network, are the strongest estimator on feature-vector (tabular) data, with the tuned neural quantile network competitive and uniquely able to sample.

Keywords: Value-at-Risk; Expected Shortfall; FRTB; quantile regression; implicit quantile networks; generative Bayesian computation; extreme value theory; conformal prediction; simulation-based calibration.

JEL classification: C14; C22; C45; C52; C58; G17; G28.

1 Introduction

Two communities measure conditional downside risk. Econometrics builds parametric filters — GARCH and its leverage and skew extensions — whose conditional scale dynamics are excellent and whose innovation law is fixed by assumption.¹ Machine learning builds distribution-free conditional quantile estimators — gradient-boosted quantile trees, neural implicit quantile networks (IQN), and, most recently, generative Bayesian computation (GBC) in the sense of Polson and Sokolov (2023) — which assume nothing about the innovation law but must learn everything from data. Industry practice, meanwhile, is neither: under the Basel Fundamental Review of the Trading Book (FRTB), the regulatory objective is Expected Shortfall at the 97.5% level, the workhorse internal models are historical simulation (HS) and GARCH-filtered historical simulation (FHS), and backtesting proceeds through VaR exception counts at 99% and 97.5%. Any claim that a new method “beats the industry standard” must therefore beat HS and FHS on $ES_{97.5}$ under exception-test discipline, not merely beat a vanilla GARCH on average loss.

This paper’s contribution is not that distribution-free methods win — unconditionally, they do not. On single-name daily equities a unit-matched GJR-GARCH- t beats our production

¹For readers outside econometrics: GARCH (generalized autoregressive conditional heteroskedasticity; Engle, 1982; Bollerslev, 1986) captures the most robust fact about markets — volatile days cluster — by letting today’s variance be a weighted average of yesterday’s variance and yesterday’s squared surprise (Eq. 2). “Leverage” extensions (GJR; Glosten et al., 1993) let down-moves raise volatility more than up-moves; “skew- t ” versions allow an asymmetric heavy-tailed innovation. What all of them fix in advance is the *shape* of the standardized surprise — the assumption this paper’s score monitors.

neural quantile network by 4.5–6.2% at every horizon (DM t between +12 and +18)²; on daily crypto, major-pair FX, and diversified equity portfolios the parametric model is as good or better. The contribution is a *predictive theory of when they win*, an engine that converts that theory into an industry-standard-and-above risk model, and a demonstration that the generative members of the family deliver two capabilities that no parametric point-forecasting model can offer at all.

1.1 The one-sentence thesis

A single measurable quantity — the local misspecification of the parametric conditional tail, operationalized as the excess kurtosis and asymmetry of recent GARCH-standardized residuals — governs when distribution-free and generative quantile methods beat industry-standard parametric VaR/ES. Where that misspecification is high they win decisively and with statistical significance; where it is low the parametric model is as good or better; and generative methods additionally deliver honest uncertainty on the risk number and likelihood-free calibration of intractable volatility models.

1.2 Contributions

1. **The misspecification frontier** (Section 4). A live, per-asset, per-day score that predicts when nonparametric quantiles beat GARCH- t , with per-date Diebold–Mariano significance, validated across four universes: CRSP US equities (top-decile edge +2.71%, DM 6.54), FX (hyperinflation tier pooled DM 4.12), twenty-six country indices (developed→frontier gradient, corr 0.53), and forty-three cross-asset instruments (corr 0.43), with Korea-2026 as a sharpening negative control: price crashes are not residual misspecification.
2. **An industry-standard-and-above risk engine** (Section 5). The residual-hybrid — GARCH conditional scale $\times$ state-conditioned nonparametric residual quantile — with an extreme-value left tail and split-conformal recalibration: co-best with CAViaR³ in the 90% Model Confidence Set⁴, significantly better than GARCH- t /FHS/EWMA/HS/GJR-skew- t , best ES_{97.5} calibration, the only entrant passing the exception tests at both levels (Kupiec+Christoffersen⁵ at 99% via the EVT tail; Kupiec at 97.5% via conformal recalibration), with the edge growing at the FRTB ten-day horizon and holding in the 2020/2022 stress window.

²Throughout, DM denotes the Diebold and Mariano (1995) test: a t -statistic on the mean difference of two forecasters’ losses (here, pinball loss), with a serial-correlation-robust standard error. Positive DM t favors the first-named model; $|t| > 1.96$ is significant at 5%.

³CAViaR (Engle and Manganelli, 2004) models the quantile itself as an autoregressive process estimated on the pinball loss — the strongest non-nested rival in the battery.

⁴The MCS (Hansen et al., 2011) is the subset of models that cannot be statistically rejected as best at a given confidence level; “co-best” means surviving the elimination sequence.

⁵Kupiec (1995) tests whether the *number* of days losses exceed VaR matches the nominal rate; Christoffersen (1998) additionally tests that exceptions arrive *independently* rather than in clusters. Both are standard in regulatory backtesting.

3. **Honest uncertainty on the risk number** (Section 6). Calibrated uncertainty bands on conditional (VaR, ES), which no standard model provides: block bootstrap on GARCH residuals (coverage 0.95 for VaR, 0.90 for 2.5% ES against a 0.90 target) completed by an extreme-value interval for the 99% ES tail (0.887). The construction is disciplined by a measured negative: a naive Gibbs posterior on raw returns understates VaR uncertainty by a factor ≈ 1.6 because serial dependence shrinks the effective sample size; residual space removes the defect, and this diagnosis — not any particular posterior — is what carries the chapter.
4. **Likelihood-free calibration of intractable models** (Section 7). Amortized GBC posteriors, validated by simulation-based calibration⁶, for Heston and rough Bergomi⁷ — non-Markovian, likelihood-intractable — recovering the Hurst roughness H with prior-uncertainty reduction 0.42 and SBC coverage at nominal. This is the capability GARCH/MLE structurally cannot supply.
5. **A generative joint-tail sampler, with its negatives** (Section 8). Direct nonparametric portfolio-tail models lose to multivariate GARCH/DCC⁸ three separate ways; a covariance-scale/nonparametric-shape hybrid restores calibration and wins the deep 1% co-crash tail, with the advantage largest on concentrated portfolios.

Magnitudes, collected in one place. Because the size of a win matters as much as its significance, we state the magnitudes up front. The *average* edge of the deployable engine over what banks run is real but modest: 0.3% of pinball over GARCH- t , 1.5% over EWMA, 1.9% over historical simulation at one day. The economically large numbers live exactly where the frontier says they should: +2.7% in the top misspecification decile (DM 6.5); 12–14% on hyperinflation FX (DM 2.9–3.8); 15% on Baltic freight and 52% on IG credit (accuracy only — calibration unresolved); a fourfold concentration of the edge from the calmest to the most turbulent volatility quintile; 1.4% at the ten-day horizon against $\sqrt{h}$ -scaling; an 8–11% ES_{97.5} over-statement by FHS/GARCH- t that the hybrid removes — the direct capital number; roughly twice the nominal rate of deep (1%) co-crash breaches under Gaussian/DCC that the hybrid restores to target; and a 42% prior-uncertainty reduction on the rough-volatility roughness parameter. Wins carry their Diebold–Mariano statistics throughout; ties and losses are quoted as ties and losses.

Two pieces of context make these magnitudes readable. First, the *asymmetry of the loss profile* is the economic argument for adoption: when the frontier is quiet, the deployable engine ties or edges the parametric benchmark (its calm-quintile edge is ≈ 0 but not negative, and by construction it nests FHS); the one configuration that loses materially — a bare standalone

⁶SBC (Talts et al., 2018): repeatedly simulate a known parameter from the prior and data from the model, fit the posterior, and check that the true parameter’s rank within its own posterior is uniform — a frequentist audit that the posterior means what it says.

⁷Rough volatility (Gatheral et al., 2018) drives variance with a fractional Brownian motion of Hurst exponent $H < \frac{1}{2}$; H measures path roughness (smaller = rougher), and the model has no closed-form likelihood.

⁸Dynamic conditional correlation (Engle, 2002): per-asset GARCH scales combined with a time-varying correlation matrix — the standard multivariate benchmark.

neural network on single names, -4.5 to -6.2% — is exactly the configuration Table 1 says never to deploy. Adopting the engine therefore risks a rounding error to collect a concentrated $+2.7\%$ in the decile where risk models actually fail, and $12\text{--}14\%$ where the parametric assumption breaks entirely. Second, for readers without a pinball intuition: a percentage pinball edge is a percentage reduction in the average forecast error at the quoted level, and its capital-relevant translation is the $8\text{--}11\%$ ES_{97.5} over-statement by FHS/GARCH that the engine removes — roughly that fraction of a desk’s market-risk capital charge held against a risk that is not there.

Throughout, the exposition follows three honesty rules fixed in advance. First, the strongest tabular estimator⁹ in the family is gradient-boosted trees, not the neural network; we present GBC as the framework, trees as the strongest estimator, and the tail-aware neural IQN as the canonical generative realization, competitive after tuning and uniquely able to sample. Second, residual kurtosis is necessary but not sufficient for a nonparametric win (carbon and corn are high-kurtosis ties; Baltic freight wins at moderate kurtosis through limit-move dynamics); the frontier is a strong regularity, not a law. Third, every negative result that shaped the design — the retired same-universe artifact, the electricity log-return artifact, the failed hierarchical blends, the failed learned gate, the three multivariate losses — is reported in the text, not the file drawer.

2 Background and industry benchmark

2.1 Risk measures and the FRTB discipline

For a return r_{t+1} with conditional law F_t , the level- α Value-at-Risk is the conditional quantile $\text{VaR}_t^\alpha = F_t^{-1}(\alpha)$ and the Expected Shortfall is $\text{ES}_t^\alpha = \mathbb{E}[r_{t+1} \mid r_{t+1} \leq \text{VaR}_t^\alpha]$. FRTB (Basel Committee on Banking Supervision, 2019) sets the capital metric at ES_{97.5} over a liquidity-adjusted horizon (ten days for the base bucket) while backtesting remains VaR-exception based: more than twelve exceptions at 99%, or thirty at 97.5%, in a 250-day year pushes a desk off the internal-models approach. The empirical bar for “beating the industry standard” is therefore joint: accuracy on ES_{97.5} and pinball, *and* unconditional (Kupiec, 1995) (the Kupiec test: are there the right number of VaR breaches?) plus independence (Christoffersen, 1998) (the Christoffersen test: are breaches independent, rather than clustered in time?) exception tests at both levels, with significance assessed by per-date Diebold–Mariano tests (Diebold and Mariano, 1995) (a t -test on the difference in loss between two forecasts) and the Model Confidence Set (Hansen et al., 2011) (the set of models statistically indistinguishable from the best one).

2.2 The parametric family, the simulation family, and CAViaR

Our battery spans the models banks actually run: historical simulation; EWMA/RiskMetrics (Zumbach, 2007) (an exponentially weighted moving average of squared returns, giving recent days more weight); GARCH(1,1)- t (Bollerslev, 1986); GJR-GARCH-skew- t (adding leverage and

⁹“Tabular” data: observations that are rows of engineered features (here, the state vector s_t) rather than raw sequences or images — the setting where tree ensembles remain the strongest general-purpose learners.

skewness; Glosten et al., 1993); GARCH-FHS (parametric scale, empirical standardized residuals — the volatile-regime best practice; Barone-Adesi et al., 1999); and SAV-CAViaR (Engle and Manganelli, 2004), included after an anti-strawman audit showed its omission flattered our model. CAViaR (conditional autoregressive Value-at-Risk) skips modeling the distribution entirely and lets the quantile itself follow a recursion estimated on the pinball loss; the SAV (symmetric absolute value) specification is $q_t(\tau) = \beta_0 + \beta_1 q_{t-1}(\tau) + \beta_2 |r_{t-1}|$ — today’s VaR is yesterday’s VaR adjusted by the size (not sign) of yesterday’s move, making it the semiparametric analogue of a GARCH recursion written directly on the quantile.

2.3 Distribution-free conditional quantiles and GBC

The nonparametric family estimates $Q_\theta(\tau \mid x_t)$ directly by pinball-loss empirical risk minimization (Koenker and Bassett, 1978): gradient-boosted quantile trees (GBM) and the implicit quantile network of Dabney et al. (2018) (IQN: the quantile level τ is itself fed into the network, so one set of weights outputs the whole quantile curve rather than one grid point at a time), the architecture at the core of GBC (Polson and Sokolov, 2023; Nareklishvili et al., 2025). Casting GBC as ERM conditional-quantile regression on observed pairs removes its forward-simulator requirement and extends it to observational time series; the generative reading — draw $\tau \sim U(0, 1)$, output $Q_\theta(\tau \mid x)$ — returns in Sections 6–8, where sampling and posterior uncertainty, not point accuracy, are the object. Generalized Bayes (a posterior built from a loss function in place of a likelihood) enters through the Gibbs posterior $\pi_n(\theta) \propto \exp\{-\omega \sum_t \ell_\theta(z_t)\} \pi(\theta)$ — a posterior built from a loss rather than a likelihood, targeting the risk minimizer directly (Jiang and Tanner, 2008; Bissiri et al., 2016), whose finance-specific failure mode and repair are the subject of Section 6; distribution-free finite-sample guarantees enter through split-conformal and adaptive conformal inference (recalibrating predicted intervals on held-out data so that coverage holds without distributional assumptions) (Gibbs and Candès, 2021).

3 Methods

This section is written to be self-contained: a reader from any one of econometrics, machine learning, or risk management should be able to reimplement every estimator from the formulas and algorithm boxes below. (For a finance-journal submission the algorithm boxes and the proof of Proposition 1 can migrate to an appendix unchanged; they are self-contained.) A glossary of all acronyms and jargon is in Appendix A.

3.1 Notation and objects

- r_t — return of an asset on day t ; $\mathcal{F}_{t-1}$ — information available before t .
- $Q_t(\tau)$ — the conditional τ -quantile of r_t given $\mathcal{F}_{t-1}$, formally $Q_t(\tau) = \inf\{q : \Pr(r_t \leq q \mid \mathcal{F}_{t-1}) \geq \tau\}$: the number the return falls below with probability τ . Value-at-Risk is its

negative, $\text{VaR}_t^\tau = -Q_t(\tau)$; Expected Shortfall is the average of the quantile curve beyond it, $\text{ES}_t^\tau = -\tau^{-1} \int_0^\tau Q_t(u) du$ — the expected loss on the days VaR is breached.

- $z_t = (r_t - \mu_t)/\sigma_t$ — the *standardized residual*: the return with the model’s conditional mean and scale divided out. If the parametric model is right, the z_t are i.i.d. draws from one fixed law.
- s_t — the state vector (trailing realized volatility, residual-shape scores, lags) on which non-parametric quantiles condition.
- $\tau \sim \lambda$ — a quantile level drawn from a base distribution λ ; $H_\phi(s, \tau)$ — a learned generator mapping (state, level) to a quantile, in the sense of Polson and Sokolov (2023): $z \stackrel{D}{=} H(s, \tau)$, $\tau \sim U(0, 1)$, is the inverse-CDF (von Neumann) representation of sampling.

Loss and metric. All quantile models are trained and scored by the pinball (check) loss

$$\rho_\tau(u) = u(\tau - \mathbb{I}\{u < 0\}), \quad u = z - q, \quad (1)$$

whose expectation is uniquely minimized by the true conditional quantile (Koenker and Bassett, 1978): differentiating $\mathbb{E} \rho_\tau(z - q)$ in q and setting it to zero gives $\Pr(z < q) = \tau$, so the loss *elicits* the quantile the way squared error elicits the mean — under-prediction is charged τ per unit, over-prediction $1 - \tau$, and the charges balance exactly at the τ -quantile. Averaging (1) over $\tau \in (0, 1)$ targets the whole distribution at once: the 1-Wasserstein distance between two laws is the L_1 distance between their quantile functions, $W_1(F, G) = \int_0^1 |F^{-1}(\tau) - G^{-1}(\tau)| d\tau$, so pinball empirical-risk minimization across levels is distribution matching in W_1 — density-free, which is what permits every estimator below to dispense with a likelihood.

Why this loss and not another. Three alternatives are standard, and the choices are deliberate. (a) The continuous ranked probability score (CRPS) is exactly the integral of the pinball loss over all levels, $\text{CRPS}(F, y) = 2 \int_0^1 \rho_\tau(y - F^{-1}(\tau)) d\tau$ (Gneiting and Raftery, 2007); scoring by CRPS therefore weights the distribution’s body and tails equally by construction. Downside risk lives at fixed regulatory levels ($\tau \in \{0.01, 0.025\}$), and the models in this paper differ almost entirely in the tail — CRPS would average those differences away against a body where all competitors agree. We therefore score pinball *at the levels that are decision-relevant*, and use tail-weighted τ -sampling in training for the same reason. (b) Expected Shortfall is not elicitable on its own — no loss function is minimized by the true ES alone — but the pair (VaR, ES) is, under the joint scoring functions of Fissler and Ziegel (2016); the FZ-scored re-evaluation of the battery is a registered extension (Section 10). (c) Likelihood is unavailable by design: several competitors (and both simulators of Section 7) have no tractable density, and likelihood scores would silently exclude them.

3.2 The final model: the residual-hybrid engine

The deployable model of this paper is *not* a bare neural network; it is a four-stage composition, each stage keeping what an industry model gets right. The design principle that survives every

experiment is a division of labor: *let the parametric filter do what it is good at (conditional scale), and spend the nonparametric capacity only on what it is bad at (the conditional shape of the innovation law)*. Intuitively: GARCH answers “how big is a typical move today?” extremely well, and badly answers “given a typical move size, how likely is a move five times that, and is down likelier than up?” The engine keeps GARCH’s answer to the first question and learns the second from data. Each finishing stage has the same one-line logic: the EVT tail says *beyond the data, extrapolate with the law extreme-value theory proves exceedances must follow, rather than trusting any fitted model outside its training range*; the conformal stage says *measure your own miss rate on data you have not touched, and shift the curve by exactly that miss* — calibration by bookkeeping, no assumptions needed.

Why one should expect this to work before seeing data. The architecture is not an empirical accident; each piece is where prior reasoning puts it. Conditional *scale* is the strongest, best-understood signal in daily returns (volatility clustering is the most robust stylized fact in the field), and it is identified by a handful of parameters — so a parametric filter estimates it with minimal variance, and a nonparametric model that re-learns it merely spends data on what four parameters already buy. Conditional *shape*, by contrast, is exactly where the parametric family is constrained by assumption and where per-asset data is too scarce to fit flexibly — tail events are rare by definition — so that is where nonparametric capacity earns its keep, and only if it can borrow strength across assets. This is classical bias–variance division of labor: parametric where the structure is known, flexible where it is not, which is why the hybrid should (and does) dominate both endpoints. Pooling residuals across names is in turn the standard hierarchical-shrinkage argument: once name-specific scale is divided out, standardized surprises are approximately draws from a common state-conditional law, and estimating a shared object from the whole cross-section improves on any single name’s estimate whenever units are scarce individually and similar collectively — the same logic that makes partial pooling dominate in hierarchical Bayes. The EVT stage is backed by a theorem rather than a hope: exceedances over a high threshold converge to the generalized Pareto family (Balkema and de Haan, 1974; Pickands, 1975), so tail extrapolation by GPD is the principled choice for *any* body model. And the conformal stage’s guarantee is Proposition 1, not an empirical observation. The experiments in this paper are therefore tests of a mechanism argued in advance, not a specification search rewarded after the fact.

Stage 1 (parametric scale). Fit GARCH- t ,

$$r_t = \mu + \sigma_t z_t, \quad \sigma_t^2 = \omega + \alpha(r_{t-1} - \mu)^2 + \beta\sigma_{t-1}^2, \quad z_t \sim t_\nu \text{ (standardized)}, \quad (2)$$

by quasi-maximum likelihood on a trailing window, and keep only $\hat{\sigma}_t$ and the residuals $\hat{z}_t = (r_t - \hat{\mu})/\hat{\sigma}_t$.

Stage 2 (nonparametric shape). Replace the fixed t_ν innovation quantile with a state-

conditioned quantile model of the residuals, trained by pinball ERM *pooled across names*,

$$\hat{\phi} = \arg \min_{\phi} \sum_{\text{names } j} \sum_{\text{days } t} \mathbb{E}_{\tau \sim \lambda} \rho_{\tau}(\hat{z}_{j,t} - q_{\phi}(\tau | s_{j,t})). \quad (3)$$

Stage 3 (EVT tail). Beyond the estimation range of (3), extrapolate with the distribution extreme-value theory says exceedances must follow (Pickands, 1975; McNeil and Frey, 2000).¹⁰ On losses $X = -\hat{z}$, fix a threshold u at the empirical $(1 - p_0)$ loss quantile with $p_0 = 0.025$, fit the generalized Pareto distribution (GPD) $G_{\xi, \beta}(x) = 1 - (1 + \xi x / \beta)^{-1/\xi}$ — shape ξ sets how heavy the tail is, scale β how wide — to exceedances $X - u > 0$ strictly on past data, and set, for $\tau \leq p_0$,

$$\hat{q}^{\text{EVT}}(\tau) = -\left[u + \frac{\beta}{\xi}((\tau/p_0)^{-\xi} - 1)\right], \quad (4)$$

which is continuous at its own endpoint, $\hat{q}^{\text{EVT}}(p_0) = -u$ (and reduces to $-[u + \beta \log(p_0/\tau)]$ in the $\xi \rightarrow 0$ exponential limit; requires $\beta > 0$ and $1 + \xi(x - u)/\beta > 0$). Since u is the *unconditional* empirical loss threshold while the Stage-2 body quantile at p_0 is the *state-conditional* $\tilde{q}(p_0 | s)$, the splice is exactly continuous only when the threshold is anchored at the body’s own p_0 -quantile, $u = -\tilde{q}(p_0 | s)$; otherwise a small level shift of size $|-u - \tilde{q}(p_0 | s)|$ remains at the join and the two pieces are spliced to agree in practice by construction. The tail is applied selectively where the raw tail under-covers (it *hurts* some volatility-index series and is not applied universally).

Stage 4 (conformal finish). On a held-out calibration split of size n , compute the signed errors $e_i = \hat{z}_i - \hat{q}(\tau | s_i)$ and shift the curve per level by the empirical correction c_{τ} chosen as the $\lceil (n+1)\tau \rceil$ -th order statistic of $\{e_i\}$: $\tilde{q}(\tau | s) = \hat{q}(\tau | s) + c_{\tau}$ (Vovk et al., 2005; Romano et al., 2019). This distribution-free layer repairs the 97.5% level that every model in the battery otherwise fails.

Proposition 1 (split-conformal validity). *If the calibration and test residuals $e_1, \dots, e_{n+1}$ are exchangeable and almost surely distinct (a continuous error distribution, or randomized tie-breaking), then for each τ with $\lceil (n+1)\tau \rceil \leq n$ the shifted quantile satisfies*

$$\tau \leq \Pr(z_{n+1} \leq \tilde{q}(\tau | s_{n+1})) = \frac{\lceil (n+1)\tau \rceil}{n+1} < \tau + \frac{1}{n+1}.$$

Proof sketch. Under exchangeability and no ties the rank of e_{n+1} among $\{e_1, \dots, e_{n+1}\}$ is uniform on $\{1, \dots, n+1\}$; the event $z_{n+1} \leq \tilde{q}(\tau | s_{n+1})$ equals $\{e_{n+1} \leq e_{(\lceil (n+1)\tau \rceil)}\} = \{\text{rank} \leq \lceil (n+1)\tau \rceil\}$, whose probability is $\lceil (n+1)\tau \rceil / (n+1) \in [\tau, \tau + \frac{1}{n+1}]$ (Vovk et al., 2005). The full

¹⁰Extreme-value theory has two classical formulations, and the distinction matters for reading the literature. The *block-maxima* approach models the maximum of each block (e.g., yearly worst loss), which converges to the generalized extreme value (GEV) family — Gumbel, Fréchet, or Weibull type depending on the tail. The *peaks-over-threshold* (POT) approach, used everywhere in this paper, models all exceedances over a high threshold, which converge to the generalized Pareto distribution (GPD) (Balkema and de Haan, 1974; Pickands, 1975). POT uses the data more efficiently (every large loss contributes, not just one per block), which is why it is the standard in risk management (McNeil and Frey, 2000; Embrechts et al., 1997). The GPD’s shape parameter ξ corresponds to the GEV type: $\xi > 0$ is the heavy-tailed (Fréchet) case relevant for financial losses.




Figure 1: The residual-hybrid engine as a pipeline. Solid arrows carry data; the dashed top arrow carries the conditional scale $\hat{\sigma}_t$ that rescales the finished residual quantile curve; dashed boxes are inputs shared across assets (the pooled panel that trains the amortized shape model) or held out from training (the calibration split behind the conformal guarantee). Every stage refits on the walk-forward schedule using only past data.

argument, the role of the no-ties hypothesis (with ties the upper bound can fail, e.g. identical residuals give coverage 1), and the boundary convention $c_\tau = +\infty$ when $\lceil (n+1)\tau \rceil = n+1$, are in Appendix B. The guarantee is marginal, not conditional on s_{n+1} . $\square$

The full forecast composes the stages:

$$\hat{Q}_t(\tau) = \hat{\sigma}_t \cdot \left[\tilde{q}_\phi(\tau | s_t) \mathbb{I}\{\tau > p_0\} + \hat{q}^{\text{EVT}}(\tau) \mathbb{I}\{\tau \leq p_0\} \right]. \quad (5)$$

Industry models as special cases. Equation (5) nests the standard toolkit. Historical simulation: $\hat{\sigma}_t \equiv 1$ and q the unconditional empirical quantile of past returns. Filtered historical simulation (FHS): q the unconditional empirical quantile of past *residuals* — Stage 2 with the state dependence deleted. GARCH- t : $q = t_\nu^{-1}$, a fixed parametric shape. The hybrid is therefore “FHS and above”: whatever FHS captures, (3) can represent, plus conditional shape. (CAViaR sits outside the nesting — it autoregresses on the quantile directly, with no scale/shape split — and is the one rival the engine only ties, Section 5.)

Algorithm 1 (residual-hybrid engine). (i) Fit (2) per asset on a trailing window; store $\hat{\sigma}_t, \hat{z}_t$. (ii) Pool $(s_{j,t}, \hat{z}_{j,t})$ across names; train the shape model by (3). (iii) Fit the GPD tail (4) on past exceedances; splice at p_0 . (iv) Learn the conformal shifts c_τ on a calibration split disjoint from the training data. (v) Forecast by (5); refit all stages on the walk-forward schedule (annual refit; all windows trailing; evaluation strictly out-of-sample).

3.3 Estimating the shape: trees for accuracy, networks for generation

Two estimators realize $q_\phi(\tau | s)$ in (3), and the paper is deliberate about which to use when (Table 1).

Gradient-boosted trees (GBM). For each level τ_k on a fixed grid, an additive ensemble of regression trees $q^{(m)}(\cdot) = q^{(m-1)}(\cdot) + \eta h_m(\cdot)$ is grown stagewise, each new tree h_m fit to the negative gradient of the pinball loss (τ_k where the current model under-predicts, $\tau_k - 1$ where it over-predicts); the fitted grid is made non-crossing by monotone rearrangement — sorting the values $\{q(\tau_1|s), \dots, q(\tau_K|s)\}$ in τ , which never increases the loss (Chernozhukov et al., 2010). On financial tabular state vectors this is the strongest *point* estimator — about 0.75% better pinball than the tuned network — consistent with the broader tabular-learning literature.

The neural implicit quantile network (IQN; the GBC estimator). A τ -conditioned network in the factorized form of Dabney et al. (2018) and Polson and Sokolov (2023),

$$H_\phi(s, \tau) = g(\psi(s) \circ \varphi(\tau)), \quad \varphi_j(\tau) = \text{ReLU}\left(\sum_{i=0}^{n_c-1} \cos(\pi i \tau) w_{ij} + b_j\right), \quad (6)$$

where $\circ$ is elementwise multiplication and g, ψ are feed-forward nets. Three modifications, each forced by a documented failure, make it competitive on financial data: tail-aware τ -sampling, $\lambda = \frac{1}{2} U(0, 1) + \frac{1}{2} \text{Beta}(0.3, 0.3)$, oversampling the extremes that uniform sampling starves (the raw net’s 99% breach rate was 3.2%); monotone rearrangement of the τ -curve; and the conformal layer of Section 3.2. The network’s earning is not point accuracy but *generation*: with $\tau \sim U(0, 1)$, $z^* = H_\phi(s, \tau)$ is an exact sampler of the fitted conditional law (the inverse-CDF representation), which is what the posterior machinery of Section 3.8 consumes and what no tree ensemble or GARCH filter provides.

How training works, concretely. The objective is the empirical version of (3): mini-batch stochastic gradient descent on $\frac{1}{|B|} \sum_{(j,t) \in B} \rho_{\tau_{jt}}(\hat{z}_{j,t} - H_\phi(s_{j,t}, \tau_{jt}))$, where each example in each batch receives a *fresh* draw $\tau_{jt} \sim \lambda$ every epoch — over training the network sees every observation at many random levels, which is what lets a single set of weights carry the entire quantile curve. Early stopping is on held-out pinball loss; monotone rearrangement and the conformal shift are applied *after* training as post-processing; and the whole pipeline (filter, shape model, tail, shifts) is refit on the annual walk-forward schedule using only past data at every step. Exact layer sizes, learning rates, and seeds are recorded in the project’s run configurations and reported with the code release rather than inline.

Amortization. One model is trained over the pooled (state, loss) pairs of hundreds of names — in generative-Bayes terms, H is an amortized map evaluated at any new (s, τ) without per-asset refitting. The intuition for why this works, and why it works at cold-start in particular, is actuarial. A per-asset GARCH learns *name-specific parameters* from that name’s own history, so with no history it has nothing. The amortized model instead learns one *universal mapping from observable state to residual shape* — “assets whose trailing volatility looks like this, with characteristics like that, produce surprises shaped like so” — estimated from millions of asset-days of *other* names. A newly listed asset is not a new problem: its first ticks and characteristics place it at a point of state space already densely populated by other names’ history, exactly as an insurer prices a new driver from the risk table built on millions of existing drivers rather than from the new driver’s empty claims file. As the name’s own history accrues, its trailing realized volatility enters s_t and the forecast tracks its own dynamics automatically — which is also the

Table 1: Which model when: the paper’s results as one decision object. Each row is established in the section cited.

Task / regime	Winner	Where
Point quantile accuracy, tabular state	GBM (trees)	§3.3
Sampling, amortized posteriors, uncertainty	Neural IQN (only entrant)	§6
Single asset, calm ($\sim 90\%$ of asset-days)	GARCH- t	§4
Single asset, top misspecification decile	Nonparametric (+2.7%, DM 6.5)	§4
Amortized cross-name panel, any regime	Nonparametric always; no gate	§4.3
Regulatory ES with a $\geq$ GARCH floor	Residual-hybrid + EVT + con-formal	§5
Multivariate portfolio tail	GARCH/DCC (honest negative)	§8

intuition for why the explicit own-history blends fail: the state vector already carries the name’s own recent information, so bolting it on again adds only noise. Feature ablation at full scale (1.32M rows, 720 names, 288 held out) shows the transfer edge is carried overwhelmingly by each name’s *own recent dynamics* (realized vol above all; removing it costs +1.0%, removing all cross-sectional characteristics only +0.5%), with characteristics mattering at cold-start. Two explicit hierarchical corrections — a naive quantile blend toward own-history and an amortized-as-prior affine Gibbs update — both *hurt* at every listing age (ratios 1.003–1.043): rich conditioning performs the partial pooling implicitly. Two external checks: on held-out names the transfer win rate over own-history benchmarks is 59–79%, and the same pipeline scores a leakage-safe weighted scaled pinball of 0.269 on the M5 benchmark — benchmark tier — outside finance altogether.

Practitioner’s selection rule. The two estimators minimize the same loss over different index sets — the trees solve $\min \sum_k \sum_i \rho_{\tau_k}(\hat{z}_i - q_k(s_i))$ on a fixed grid $\{\tau_k\}$, one model per level; the network solves $\min_{\phi} \sum_i \mathbb{E}_{\tau \sim \lambda} \rho_{\tau}(\hat{z}_i - H_{\phi}(s_i, \tau))$, one model for the continuum — and the choice follows from what the application consumes. If the desk needs *numbers at fixed regulatory levels* (VaR/ES reporting, limits, capital), use the trees: best point accuracy, no GPU, no sampling needed. If the application consumes *draws* — scenario generation, portfolio aggregation by simulation, the posterior and SBC machinery of Sections 6–7 — use the network: it is the only member that is an exact sampler, at a measured cost of $\sim 0.75\%$ in point accuracy. If both are needed, run both from the same pooled residual panel; they share every upstream and downstream stage of (5), so the swap is one component, not a new pipeline.

Algorithm 2 (amortized shape model). (i) Assemble the pooled panel $\{(s_{j,t}, \hat{z}_{j,t})\}$ over names j and days t . (ii) Train H_{ϕ} by SGD on $\mathbb{E}_{\tau \sim \lambda} \rho_{\tau}(\hat{z} - H_{\phi}(s, \tau))$ [network], or boost per- τ regressions on a grid [trees]; enforce monotonicity by rearrangement. (iii) For any asset — including one with no return history, using characteristics-only s — evaluate $H_{\phi}(s_t, \cdot)$; no per-asset estimation step exists. (iv) To *sample*, draw $\tau^{(m)} \sim U(0, 1)$ and return $z^{(m)} = H_{\phi}(s_t, \tau^{(m)})$.

3.4 Formal results: why shape misspecification must show up in the loss

The paper’s organizing claim — that the nonparametric edge lives exactly where the fixed innovation shape is locally wrong — is not only an empirical pattern; it follows from the structure of the pinball loss. The results below are elementary but do real work: they convert “the model’s shape assumption is violated” into a quantitative lower bound on forecastable loss.

Lemma 1 (pinball regret identity). *Let $Z \sim F$ with $\mathbb{E}|Z| < \infty$ (so S is finite) and $q_F = F^{-1}(\tau)$, and let $S(q) = \mathbb{E}_F \rho_\tau(Z - q)$ be the expected pinball loss of forecasting the constant q . Then for any q ,*

$$S(q) - S(q_F) = \int_{q_F}^q (F(u) - \tau) du \geq 0, \quad (7)$$

with equality iff $F \equiv \tau$ on the interval between q_F and q . If F has density bounded as $0 < m \leq f \leq M$ on that interval,

$$\frac{m}{2} (q - q_F)^2 \leq S(q) - S(q_F) \leq \frac{M}{2} (q - q_F)^2. \quad (8)$$

Proof sketch. The difference quotients of $q \mapsto \rho_\tau(Z - q)$ are dominated by the Lipschitz constant $\max(\tau, 1 - \tau)$, so differentiation under the expectation is justified and $S'(q) = F(q^-) - \tau$ (a.e. equal to $F(q) - \tau$); integrating from q_F to q gives (7) for any F , atoms forming a null set. The integrand $F(u) - \tau$ has the sign of $u - q_F$ (by definition of q_F), so the integral is ≥ 0 whichever side q lies; (8) then follows from $F(q_F) = \tau$ (supplied by the density hypothesis) and $m|u - q_F| \leq F(u) - \tau \leq M|u - q_F|$. Full details, including both orientations of the integral, are in Appendix B. $\square$

Lemma 1 says the extra loss a misspecified forecaster pays is *quadratic in its quantile error*. The frontier claim then reduces to: does shape misspecification create a quantile wedge at the risk levels that matter? For the paper’s canonical case — a fixed Student- t shape when the true local shape is fatter — it must:

Proposition 2 (tail wedge under kurtosis misspecification). *Let $Q_\nu(\tau)$ denote the τ -quantile of the unit-variance Student- t with $\nu > 2$ degrees of freedom. There exists $\tau^*(\nu_1, \nu_2) \in (0, 1/2)$ such that for all $\tau < \tau^*$ and $\nu_1 < \nu_2$, $Q_{\nu_1}(\tau) < Q_{\nu_2}(\tau)$: the fatter-tailed law has the strictly more extreme tail quantile, and the wedge $|Q_{\nu_1}(\tau) - Q_{\nu_2}(\tau)| \rightarrow \infty$ as $\tau \downarrow 0$. Consequently, by Lemma 1, a forecaster committed to shape t_{ν_2} while the local truth is t_{ν_1} pays regret bounded below by $\frac{m}{2} (Q_{\nu_1}(\tau) - Q_{\nu_2}(\tau))^2$ at every level $\tau < \tau^*$, where m is the minimum of the true density on the wedge interval.*

Proof sketch. The unit-variance t_ν ($\nu > 2$, so the variance is finite and the normalization is well defined) has a regularly varying left tail: $\Pr(Z < -x) = c_\nu x^{-\nu}(1 + o(1))$ as $x \rightarrow \infty$ with $c_\nu > 0$; unit-variance rescaling alters c_ν but not the index ν . Inverting a regularly varying tail (Embrechts et al., 1997) gives $Q_\nu(\tau) = -(c_\nu/\tau)^{1/\nu}(1 + o(1))$ as $\tau \downarrow 0$. For $\nu_1 < \nu_2$ the exponent $1/\nu_1 > 1/\nu_2$ makes $|Q_{\nu_1}(\tau)|/|Q_{\nu_2}(\tau)| \rightarrow \infty$, so the strict ordering and the diverging wedge hold for all τ below some $\delta > 0$; taking $\tau^* = \min(\delta, 1/2)$ places τ^* in $(0, 1/2)$ as claimed. (The

proposition asserts *existence* of such a τ^* , not the location of the finite-sample crossover; the latter is $\tau_c \approx 0.0179$ for the t_3 -vs-normal case of Remark 1.) The regret bound is (8); full details in Appendix B. $\square$

Remark 1 (the crossover is real and matters). Because unit-variance fat-tailed laws put *more* mass near the center as well as in the far tail, the quantile ordering of Proposition 2 reverses at moderate levels: numerically, the unit-variance t_3 has a *less* extreme 5% quantile than the normal (-1.36 vs -1.64) but a more extreme 1% quantile (-2.62 vs -2.33). The sign flip occurs at a single crossover level $\tau_c \approx 0.0179$ (solving $Q_{t_3}(\tau) = \Phi^{-1}(\tau)$ numerically for the lower-tail root), so the deepest deployable level $\tau = 0.01$ sits *below* the crossover (fat tail strictly more extreme) while $\tau = 0.025$ sits just *above* it (at 2.5% the unit-variance t_3 quantile is -1.837 , still less extreme than the normal’s -1.960). Two of the paper’s design choices are direct consequences. First, this is precisely why CRPS-style whole-distribution scoring can rank a fat-tail-aware model *below* a thin-tailed one even when the tail is the decision-relevant region (Section 3.1): body and far tail disagree in sign. Second, it is why the deepest deployable level straddles the crossover rather than sitting cleanly beyond it, and why the score of Section 3.7 is used as a *rank* rather than a moment estimate.

Lemma 2 (validity of the generative sampler). *If $\tau \mapsto \tilde{q}(\tau \mid s)$ is nondecreasing and left-continuous (guaranteed by the rearrangement step) and $\tau \sim U(0,1)$, then $Z^* = \tilde{q}(\tau \mid s)$ has quantile function exactly $\tilde{q}(\cdot \mid s)$.*

Proof sketch. For any x , left-continuity and monotonicity make $\{t \in (0,1) : \tilde{q}(t \mid s) \leq x\}$ an interval closed at its top, so $\{\tilde{q}(\tau \mid s) \leq x\} = \{\tau \leq G(x)\}$ with $G(x) = \sup\{t : \tilde{q}(t \mid s) \leq x\}$. Since $\tau \sim U(0,1)$, $\Pr(Z^* \leq x) = \Pr(\tau \leq G(x)) = G(x)$, i.e. Z^* has CDF G . As G is the generalized inverse of $\tilde{q}$ and $\tilde{q}$ is nondecreasing and left-continuous, the generalized inverse of G returns $\tilde{q}$ exactly (the standard inverse-of-inverse duality), so $\tilde{q}(\cdot \mid s)$ is the quantile function of Z^* . Full statement in Appendix B. $\square$

Remark 2 (nesting and ERM dominance). The hypothesis class of Stage 2 contains the constant-in- s map $q(\tau)$, i.e., FHS; empirical risk minimization over the larger class therefore attains in-sample pinball risk no worse than FHS by construction, and the out-of-sample comparisons of Section 5 test whether that dominance survives estimation noise. It does; the honest converse is that the same argument does *not* apply to CAViaR, which is why CAViaR is the one rival the engine only ties.

3.5 A synthetic boundary check: when the frontier does *not* appear

Before the real data, a simulation with known truth — run once, seed fixed — checks the formal results and, more usefully, delineates when the frontier edge should *fail* to appear. The data-generating process is deliberately favorable to the thermometer: unit-variance Student- t residuals whose degrees of freedom switch between a calm regime ($\nu = 12$) and a fat regime ($\nu = 3$) under a persistent Markov chain; forecasters see the trailing 63-day kurtosis score and

compete at $\tau \in \{0.01, 0.025\}$ (80,000 days; scale known, isolating shape). Three results. First, the Lemma 1 identity verifies numerically: the regret integral and the Monte-Carlo regret agree to five decimals (3.5×10^{-5} at $\tau = 0.01$). Second, the conformal shift moves realized coverage toward nominal on the held-out split, as Proposition 1 promises. Third — the instructive one — the score-binned nonparametric forecaster does *not* beat the fixed shape here: a single $t_{\hat{\nu}}$ fitted on the mixture ($\hat{\nu} = 6.8$) is nearly unbeatable at these levels, with total regret under 0.5% and no decile gradient. Two mechanisms explain it, and both teach. (i) By Lemma 1 regret is *quadratic* in the quantile wedge, and at $\tau \geq 0.01$ the wedge between the fitted compromise and either regime is small — static kurtosis *blending* is a second-order sin. (ii) The sample kurtosis of a t_{ν} law has infinite estimator variance for $\nu \leq 8$, so a moment-based thermometer is at its noisiest precisely when tails are fat, and sixty-three observations cannot reliably rank regimes that differ only in a static ν . The lesson for reading the empirical sections is sharp: the real-data top-decile edge of Section 4 cannot be an artifact of stationary fat tails being averaged over — that configuration is exactly what the simulation shows a fixed shape handles — but must reflect *local* shape breaks a single fitted ν cannot straddle: asymmetry, jump clustering, and residual misfit that co-move with the score. This is the same conclusion the Korea 2026 negative control reaches from the opposite direction (price crashes are not residual misspecification), and it is why the score combines kurtosis with asymmetry and a jump indicator rather than kurtosis alone. (Simulation script `toy_example.py` ships with the code release.)

3.6 What is different from Polson–Sokolov GBC

Because Polson and Sokolov (2023) is this paper’s methodological benchmark, we state the deltas explicitly; each is an innovation forced by the downside-risk target, and each is defined mathematically above. (1) *Base measure*: GBC trains with $\tau \sim U(0, 1)$; we train with the tail-weighted mixture $\lambda = \frac{1}{2}U(0, 1) + \frac{1}{2}\text{Beta}(0.3, 0.3)$, because uniform sampling starves the extreme levels that downside risk is about (the uniform-trained net’s 99% breach rate was 3.2%). (2) *Input space*: GBC maps raw data to posteriors; we first pass returns through the parametric filter (2) and model *standardized residuals* — the scale dynamics are the one thing the industry model already gets right, and Section 6 measures the cost of skipping this step (a 1.6× overconfident posterior). (3) *Beyond-data tail*: no generative regression can know about quantiles beyond its training support; we splice the GPD tail (4) at p_0 , replacing extrapolation by the network with extrapolation by extreme-value theory. (4) *Finite-sample guarantee*: GBC’s output carries no coverage guarantee; the conformal stage adds one (Proposition 1). (5) *Amortization target*: GBC amortizes across *simulations of one model*; we amortize across the real cross-section of hundreds of assets, which is what creates the cold-start capability and the transfer results of Section 3.3. (6) *Scope discipline*: GBC is a computational method; this paper adds the empirical question GBC does not ask — *when* such machinery beats the parametric filter at all — answered by the score below.

3.7 The misspecification score

GARCH- t assumes standardized residuals $z_t = (r_t - \mu_t)/\sigma_t$ are i.i.d. Student- t with fixed shape. Static heavy tails do not violate this — ν absorbs them. What violates it is the shape being *locally* wrong: recent residuals far fatter-tailed or more asymmetric than any fixed t . We therefore score, per asset-day, over the trailing window $W_t = \{t - 62, \dots, t\}$ of $n_w = 63$ residuals with window mean $\bar{z}$ and standard deviation s_z :

$$\text{mk}_{63}(t) = \frac{1}{n_w} \sum_{i \in W_t} \left(\frac{\hat{z}_i - \bar{z}}{s_z} \right)^4 - 3, \quad \text{sk}_{63}(t) = \frac{1}{n_w} \sum_{i \in W_t} \left(\frac{\hat{z}_i - \bar{z}}{s_z} \right)^3, \quad (9)$$

Terminology is fixed here once: *skewness* is the signed third moment sk_{63} ; *asymmetry* always means its absolute value $|\text{sk}_{63}|$ (departure from symmetry in either direction, since a fixed symmetric t is violated equally by either sign); everywhere the text says “asymmetry,” the statistic computed is $|\text{sk}_{63}|$. The score also uses the five-day jump indicator $J_5(t) = \max_{i \in \{t-4, \dots, t\}} |\hat{z}_i|$. The frontier score is the decile rank of these within the trailing cross-section. All three are causal (trailing-window) and computable on any asset with a fitted GARCH — which is what makes the frontier a deployable monitor rather than an ex-post classification. Informally we will call this the misspecification *thermometer*: it reads the local temperature of the parametric assumption, and Table 1 is the treatment chart.

Why these statistics. The parametric model’s one fragile assumption is that the standardized residual’s *shape* is a fixed t_ν ; its first two conditional moments are already absorbed by μ_t, σ_t . The lowest-order ways a fixed symmetric law can be locally wrong are precisely the third and fourth standardized moments: sk_{63} detects asymmetry no symmetric t can carry, and mk_{63} detects tail mass beyond what the fitted ν implies (an exact t_ν sample has known positive excess kurtosis; the score is used as a *rank*, so only its cross-sectional and temporal variation matters, not its t_ν baseline). J_5 catches jump episodes too recent for the window moments to register. When these are elevated, the ERM shape model (3) has something to learn that t_ν^{-1} cannot represent — and Section 4 shows the realized edge concentrates exactly there. Note the logical status of this choice: after μ_t and σ_t are fitted, the fixed innovation shape is the parametric model’s *only* remaining free assumption. A meter that measures the local violation of exactly that assumption is therefore not one candidate signal among many — it is the residual degree of freedom itself, which is why one should expect it, a priori, to be the axis along which flexible-shape models win.

Algorithm 4 (the misspecification meter, as deployed). (i) Fit (2) on a trailing window; store $\hat{z}_{t-62}, \dots, \hat{z}_t$. (ii) Compute mk_{63} , $|\text{sk}_{63}|$, J_5 by (9). (iii) Convert each to a decile rank against the trailing pooled panel (same asset class); the score is the rank of the maximum. (iv) Act by Table 1: top decile $\Rightarrow$ the standard parametric VaR is currently, locally wrong by $\sim 2\text{--}3\%$ of pinball with high confidence — flag the asset for the nonparametric / residual-hybrid engine and for model-risk review; elsewhere $\Rightarrow$ no action (in the amortized setting the engine is run regardless; the meter is surveillance, not a switch, Section 4.3). (v) Latency: detection lags episode onset by median 0 / mean 2–4 days; the edge is back-loaded within episodes, so the lag forfeits little.

3.8 Generative posterior machinery

For Section 6 we form (i) Gibbs (generalized-Bayes) posteriors over the level- α quantile q : with prior π , learning rate ω , and the loss itself in the exponent,

$$\pi_\omega(q \mid \hat{z}_{1:n}) \propto \pi(q) \exp \left\{ -\omega \sum_{t=1}^n \rho_\alpha(\hat{z}_t - q) \right\}, \quad (10)$$

with ω either naive ($\omega = 1$) or calibrated by matching the posterior SD to a moving-block bootstrap SD. An honest caveat governs how (10) may be used. The Gibbs posterior was built to target the *risk minimizer*: its mode estimates the quantile consistently for *any* ω , but the spread of the posterior — and hence the width of any credible interval — scales with ω , and nothing in the construction guarantees those intervals cover at their stated rate. This is a known problem with a known treatment: Syring and Martin (2019a) tune ω by direct calibration — iterate until bootstrap-estimated coverage of the credible sets hits nominal (the GPC algorithm) — and applied exactly this program to Value-at-Risk (Syring and Martin, 2019b). Our SD-matching rule is a one-step, moment-based cousin of GPC: it matches the posterior’s second moment to a block-bootstrap estimate of the estimator’s true sampling variability, which coincides with coverage calibration when the posterior is approximately Gaussian but carries no guarantee beyond that; the full GPC iteration on dependent data (block-bootstrap inner loop) is a registered extension. Concretely, ω is tuned as follows. *Why ω matters*: multiplying the loss by ω in (10) is equivalent to pretending the sample carries ω -times its actual information, so large ω shrinks the posterior (overconfidence) and small ω inflates it. *Our one-step rule*: (i) compute the block-bootstrap standard deviation $\hat{se}_{\text{boot}}$ of the point estimate — an assumption-light estimate of its true sampling variability under dependence; (ii) compute the posterior SD at $\omega = 1$; (iii) rescale $\omega \leftarrow \omega \cdot (\text{SD}_{\text{post}} / \hat{se}_{\text{boot}})^2$, using the fact that posterior variance scales as $1/\omega$; one step lands the posterior SD on the bootstrap SD. *Why this is reasonable*: for an approximately Gaussian posterior, an interval with the right SD and the right center has the right coverage; GPC generalizes exactly this by iterating on coverage itself rather than the SD proxy. This is also why the chapter’s deployable intervals are the bootstrap and EVT constructions, with the calibrated Gibbs posterior as a supporting variant: its principal contribution here is diagnostic — the $1.6\times$ raw-return overconfidence that residual space repairs — rather than the intervals

themselves. Two design rules are load-bearing, each confirmed by an ablation. The exponent must be the *pinball* loss: an expectile (squared-asymmetric) exponent mis-centers the posterior and collapses coverage to 0.1–0.4. And (10) must be formed on GARCH *residuals*, not raw returns: volatility clustering makes the effective sample size of raw returns far smaller than n , and the naive raw-return posterior is $1.6\times$ overconfident on VaR. We also form (ii) block and i.i.d. bootstrap distributions of $(\widehat{\text{VaR}}, \widehat{\text{ES}})$, and (iii) a GPD peaks-over-threshold interval for deep-tail ES.

For Section 7 we use the standard generative-Bayesian computation (GBC) recipe of Polson and Sokolov (2023), whose validity check is simulation-based calibration (SBC) (Talts et al., 2018).

Algorithm 3 (likelihood-free posterior with SBC audit). (i) Simulate $\theta^{(i)} \sim \pi(\theta)$ and paths $y^{(i)} = f(\theta^{(i)})$ from the forward model (Heston, rough-Bergomi) — no likelihood is evaluated anywhere. (ii) Reduce each path to summaries $S(y^{(i)})$ (realized-variance moments, autocorrelations, tail counts). (iii) Train amortized posterior quantiles $Q_\phi(\tau \mid S)$ for each component of θ by pinball ERM (IQN and GBM). (iv) *Audit*: for fresh simulations with known θ^{true} , compute the rank of θ^{true} in its own posterior; ranks must be uniform and stated intervals must cover at the advertised rate (SBC). (v) Only then evaluate at the observed data’s summaries $S(y)$.

A natural objection to step (i) deserves a direct answer: *if you simulate the training data yourself, you know the true parameters — why bother with a posterior at all?* Because the simulations are not the object of interest; the *map* is. Each simulated pair $(\theta^{(i)}, y^{(i)})$ is one example of “data that a model with known parameters produces”; training on many such pairs, drawn across the whole prior, teaches Q_ϕ the inverse relationship — from any dataset to the distribution of parameters that could have generated it. The known truths are spent on *learning and auditing* that map (steps iii–iv), never on the real data; the map’s value is at step (v), where it is evaluated on the one dataset whose parameters nobody knows. The situation is the same as training a translator on millions of sentence pairs whose translations are known, precisely so it can translate the sentence whose translation is not.

4 The misspecification frontier

How to read every number in this and later sections. Conventions are fixed once, here.

(i) *Edges* are percentage reductions in pinball loss, so *larger is better* for the named model; a “+2.7% edge” means the model’s average loss is 2.7% lower than the benchmark’s. (ii) *Ratios* are model-over-benchmark loss ratios, so *below 1 favors the model* and $1 - \text{ratio}$ is the edge ($0.973 =$ a 2.7% win; $1.043 =$ a 4.3% loss). (iii) *DM* is the Diebold–Mariano t -statistic on the loss difference: its sign follows the first-named model, and $|t| > 1.96$ is significance at the 5% level (> 2.58 at 1%). (iv) *Coverage* numbers are realized frequencies against a stated nominal (0.05 target: 0.056 is slight under-coverage of risk, 0.032 over-conservatism); for *credible-interval* coverage, closer to nominal is better in both directions. (v) *Correlations* quoted for the frontier relate a universe’s

Table 2: The misspecification frontier: nonparametric edge over GARCH- t by decile of the trailing residual-shape score (200 CRSP names, 221.6k test asset-days). Edge = (GARCH pinball – nonparam pinball)/GARCH pinball.

Decile	Edge by signal decile (%)		
	mk ₆₃ (kurtosis)	sk ₆₃ (asymmetry)	jump ₅
1–8 (bulk)	≈ 0	≈ 0	mixed, ≈ 0
9	+0.52	+0.69	—
10 (top)	+2.46	+2.37	+0.90
Overall	+0.44		

Per-date Diebold–Mariano: top mk₆₃ decile edge +2.71%, DM 6.54 ($p \approx 0$); bottom decile +0.19%, DM 1.41 ($p = 0.079$, not significant); top two deciles pooled +1.65%, DM 6.74; overall +0.44%, DM 5.67. Top-decile signal means: kurtosis ≈ 25 , |skew| ≈ 4 .




Figure 2: The misspecification frontier. Nonparametric pinball edge over GARCH- t by decile of the trailing 63-day residual excess kurtosis (left) and residual asymmetry (right), 200 CRSP names, 221.6k test asset-days. Deciles 1–8 are ≈ 0 (drawn at zero; exact bulk values in the result file); the edge jumps in the top decile (+2.46% kurtosis, +2.37% asymmetry; per-date DM 6.5 in the top kurtosis decile).

residual-kurtosis level to the size of the nonparametric edge across that universe’s members. No other significance marks are used.

4.1 The within-universe result

On 200 CRSP names (≈ 221.6 k test asset-days), we bucket the per-day pinball edge of the amortized nonparametric model over GARCH- t into deciles of each misspecification signal. Table 2 shows the result that organizes the paper.

The edge is a top-decile phenomenon. It is large and significant exactly where the fixed- t shape assumption is currently, locally broken, and statistically indistinguishable from zero where it is not. The low overall correlation between signal and edge (0.05–0.10) is not a weakness of the story but its point: the relation is a threshold nonlinearity, not a linear trend, and the top-decile Diebold–Mariano statistics are the correct summary.

Table 3: The frontier across universes. Ratio = nonparam/GARCH- t pinball (< 1 : nonparametric wins). All fits stand

Universe	Segment	Ratio	Resid. kurt.	S
FX (24 pairs)	majors (7)	1.019	39	G
	EM (10)	1.013	8	G
	crisis (7)	0.973	~ 1915	n
	USDARS	0.880	3298	D
	USDEGP	0.857	1920	D
Country indices (26)	developed (8)	1.007	4.7	0
	emerging (10)	1.012	5.3	0
	frontier (8)	0.993	10.0	4
Cross-asset (43)	vol indices	~ 0.987	32.8	3
	Baltic freight	0.846	—	D
	IG credit OAS	0.960	—	D
	13/43 significant wins overall; corr(log kurt, edge) = 0.43			
Korea 2026 (23)	all types	1.01–1.12	4–13	G

^a Sri Lanka (DM 3.20), Kenya (3.52), Pakistan (2.38), Nigeria (2.06). Cross-index corr(log residual kurtosis, edge) = 0.53. Phi despite kurtosis 13.6), illustrating necessity-not-sufficiency.

4.2 The frontier as geography: four universes

Table 3 shows the same axis at work across markets. Three features deserve emphasis. First, the only decisive *standalone* single-asset wins anywhere are the hyperinflation and capital-control currencies — Argentine and Egyptian pesos, 12–14% better with significance — whose residual kurtosis (~ 2000 – 3300) sits two orders of magnitude above anything GARCH- t was designed to absorb. Second, the country gradient is mediated by residual kurtosis, not by the development label per se: frontier indices (kurt ~ 10) sit between well-behaved developed indices (kurt ~ 5) and hyperinflation FX on the same frontier. Third, the Korea-2026 negative control sharpens the mechanism: a dramatic *price* crash with circuit breakers is *not* residual misspecification — Korean residual kurtosis stays in the ordinary 4–13 range, and GARCH- t significantly wins all twenty-three assets even inside the crisis window. The frontier is about the residual shape, not about headline turbulence.

A frequency corollary. Daily crypto belongs to GARCH (zero nonparametric wins on five liquid coins). At *hourly* frequency the picture inverts: the neural IQN beats the full benchmark ladder, including filtered historical simulation, by $\sim 0.6\%$ with per-date DM $p < 10^{-7}$. Sample density is itself a frontier dimension: more observations per unit of regime raise the resolvable residual-shape signal. A systematic intraday map of the frontier is registered and blocked only on tick-data access.




Figure 3: The frontier as geography: nonparametric edge vs post-GARCH residual kurtosis, tier means and single assets with study-reported kurtosis only (country universe corr 0.53; cross-asset corr 0.43). The hyperinflation-FX tier (kurtosis ~ 2000 – 3300) carries the only decisive standalone wins; high kurtosis is necessary but not sufficient — carbon (288) and corn (20) tie, and NGN’s win is not significant. The Korea-2026 control (kurtosis 4–13, GARCH wins all 23 assets) sits in the lower-left regime.

4.3 Regime concentration and the futility of gating

In the amortized panel the edge concentrates in turbulence: the GARCH-minus-GBM pinball gap grows monotonically from 0.0059 in the calmest realized-vol quintile to 0.0247 in the most turbulent — roughly fourfold — while both nonparametric estimators beat GARCH in *every* quintile (overall: GARCH 0.6543, GBM 0.6447, IQN 0.6478). A natural inference is that one should *gate*: run GARCH on calm days and switch to the nonparametric model when the score is high. It fails, instructively. A classifier gate trained on the misspecification signals routes 90.5% of days to the nonparametric model and exactly matches always-nonparametric (0.3987 vs GARCH 0.4000); worse, its routing runs *backwards* across score deciles, because nonparametric wins on high-score days are rare-but-large while a classifier targets frequency. A regression gate on the edge magnitude fixes the direction (routing flat ≈ 1.0 , predicted-edge correlation +0.16) and still converges to always-nonparametric. The per-day oracle gap (4.4%) is unreachable because per-day model advantage is dominated by *which day the tail event lands* — unpredictable from prior-day state. The deployable conclusions: in the amortized setting, use the nonparametric model always (it does not underperform on calm days); use the residual-hybrid where a guaranteed $\geq$ GARCH floor is wanted; and use the frontier score as a *monitor* (Section 9), not a switch. (The score’s detection latency is small — median zero, mean two to four days — and because the nonparametric edge is back-loaded within misspecification episodes, the lag costs little even in the per-name setting, where gating does matter because standalone nonparametric models underperform on calm days.)

Table 4: FRTB battery: average pinball loss (7 quantile levels), $ES_{97.5}$ calibration, and 99% exception tests. 140 CRSP names, 155k test observations.

Model	Avg. pinball	$ES_{97.5}$ pred. vs realized	Breach ₉₉	Kupiec ₉₉ p
<i>Ideal value</i>	<i>(smaller better)</i>	<i>(pred. = realized)</i>	<i>1.00%</i>	<i>> 0.05 (pass)</i>
Residual-hybrid (GBM)	0.3551	−5.80 / −5.76	1.19%	0.00
+ EVT tail	0.3554	−6.08 / −5.86	1.04%	0.10
GARCH(1,1)- t	0.3561	−6.17 / −5.65	1.06%	0.026
GARCH-FHS	0.3565	−6.20 / −5.59	1.05%	0.041
GJR-GARCH-skew- t	0.3574	—	—	—
EWMA (RiskMetrics)	0.3606	—	—	—
Historical simulation	0.3619	—	—	—

Diebold–Mariano, hybrid vs: GARCH- t 4.38, FHS 5.70, GJR 5.73, HS 7.34, EWMA 8.33 (all $p \approx 0$). Hybrid+EVT still beats GARCH- t (DM 3.90) and FHS (DM 4.28); its 0.0003 pinball concession to the raw hybrid (DM −5.06) buys the exception-test pass. Hybrid+EVT passes Christoffersen₉₉ ($p = 0.20$); no other model in the battery passes both 99% tests while topping the accuracy table. At 97.5% *every* raw model fails Kupiec; the split-conformal recalibration repairs it (GBM-recal $p = 0.24$; IQN-v2-recal $p = 0.15$, Christoffersen $p = 0.068$).

Reading guide: Breach₉₉ is the realized frequency of losses beyond the 99% VaR — ideal is exactly 1.00% (higher = risk understated, lower = capital wasted); the ES column compares predicted to realized average exceedance loss, ideal is equality; Kupiec p above 0.05 means the breach count is statistically consistent with 1%. Bold marks the best value in each column. “—” entries were not recorded in the archived runs for the lower-accuracy tier (their exception statistics did not affect any ranking); completing them is scheduled with the registered battery re-run.

5 Beating the industry standard: the FRTB battery

5.1 Head-to-head accuracy and significance

Table 4 reports the full battery on 140 CRSP names (155k obs, 1108 dates).

Three honesty items qualify the headline. First, adding SAV-CAViaR — absent from banks’ standard toolkit but the strongest semiparametric academic benchmark — produces a statistical tie: CAViaR 0.3461 vs hybrid 0.3460 (DM 0.59, $p = 0.28$); the 90% Model Confidence Set contains exactly these two. The hybrid is *co-best with CAViaR*, not uniquely dominant, and still significantly beats everything banks actually run. Second, the neural IQN as residual model initially loses to trees badly (0.3650 vs 0.3551, DM 13.3) with a severely thin tail (breach₉₉ 3.2%); the tail-aware, monotone, conformally recalibrated version closes to a hair (0.3638 vs 0.3632, DM 3.03) with calibration repaired (breach₉₉ 0.92%). Trees remain the strongest tabular estimator; the tuned IQN is competitive and is the estimator that generalizes to Sections 6–8. Third, $ES_{97.5}$ calibration — the actual capital metric — is where the hybrid’s advantage is most consequential: FHS and GARCH- t both *over-state* ES by 8–11%, i.e. over-charge capital, while the hybrid is nearly exact.




Figure 4: The FRTB battery. Left: average pinball loss by model (axis truncated to resolve the ordering; differences carry DM 4.4–8.3 vs the hybrid, all $p \approx 0$); CAViaR, not shown, statistically ties the hybrid. Right: $ES_{97.5}$ predicted vs realized — FHS and GARCH- t over-state ES by 8–11% (over-charging capital) while the hybrid is nearly exact.




Figure 5: The edge grows with horizon. Nonparametric/GARCH pinball ratio at $h = 1, 5, 10, 20$ days (amortized panel): fat tails compound while $\sqrt{h}$ scaling compounds the shape error — $\sim 5\%$ at $h = 20$.

5.2 The ten-day horizon and stress

At the FRTB ten-day liquidity horizon (CRSP 2014–2024, directly-learned ten-day quantiles vs GARCH $\sqrt{h}$ -scaling) the pinball edge grows from $\sim 0.4\%$ at $h = 1$ to $\sim 1.4\%$ at $h = 10$, and holds in the 2020–2022 stress window (hybrid $1.247 < \text{FHS } 1.2605 < \text{GARCH } 1.262$).

More consequentially for capital: ten-day GARCH $\sqrt{h}$ -scaled ES is badly over-stated (predicted -19.4 vs realized -15.4) while the directly-learned hybrid is calibrated (-19.05 vs -17.70) with the best ten-day breach $_{97.5}$ (0.0265 vs target 0.025 ; GARCH 0.033). Fat tails compound with horizon; $\sqrt{h}$ Gaussian-style scaling compounds the shape error. A caveat: the temporal split places the whole test window in 2020–2024; a calm-vs-stress comparison awaits the 2000–2024 WRDS panel (registered, Section 10).

5.3 Cross-asset winners, with corrections

The clean calibrated cross-asset win is VIX (pinball 1.127 vs 1.137, DM 1.74, $p = 0.041$, passes Christoffersen at both levels). Credit (52% accuracy edge, DM 30.6) and Baltic freight (15%, DM 9.4) fail breach-independence — exceptions cluster — and are reported as accuracy edges with open calibration problems, not wins. An earlier electricity “win” was an artifact: log-returns on near-zero and negative power prices; refit on price *differences*, GARCH wins ERCOT (DM -2.55) and PJM (DM -2.73), and electricity is removed from the win list. The GPD tail also *hurts* some volatility series (VVIX, V2X, MOVE flip from win to tie/loss), motivating the selective-application rule of Section 3.2.

6 Honest uncertainty on the risk number

Every model so far outputs a *point* VaR/ES. None tells the desk how uncertain that number is. This section delivers calibrated uncertainty bands on (VaR, ES) and, on the way, settles a methodological question: whether the Gibbs-posterior program of generalized Bayes (Jiang and Tanner, 2008; Bissiri et al., 2016) — a posterior over the risk functional built from the loss itself — is usable on financial data at all. Its guarantees assume near-independence; financial returns are anything but. The answer, in numbers: not on raw returns (a factor-1.6 overconfidence we measure directly), yes after GARCH filtering — but the workhorses of the calibrated construction turn out to be the block bootstrap (resampling contiguous blocks of days, not single days, so the resample preserves serial dependence) (Künsch, 1989), adaptive conformal inference¹¹, and an extreme-value tail interval, with the repaired Gibbs posterior as a supporting variant rather than the centerpiece. One further generalized-Bayes result deserves record: in the project’s pre-registered daily single-portfolio study — the honest anchor in which a leverage-aware GARCH was the best standalone model and the standalone network only reached parity — a Gibbs-posterior model *average* that included the network did best of all. Loss-based posterior weighting earns its keep as an ensembling device even where the standalone generative model does not win.

6.1 The defect, measured

On 80 CRSP names at $\tau = 0.05$, we compare the Gibbs posterior SD of the VaR level to the true sampling SD estimated by moving-block bootstrap (block 20). On *raw returns*: Gibbs SD 0.108 vs block SD 0.179 — the naive posterior understates VaR uncertainty by a factor $R = 1.64$ (median 1.59). The i.i.d.-bootstrap ratio is 1.19, so the 1.19→1.64 gap isolates the serial-dependence contribution: vol clustering makes the effective sample size far smaller than n , and the product-form posterior over-concentrates accordingly.

¹¹ACI (Gibbs and Candès, 2021) treats miscoverage as a control problem: after each day the working level is nudged, $\alpha_{t+1} = \alpha_t + \gamma(\alpha - \mathbb{I}\{\text{breach at } t\})$, so a run of breaches automatically widens the band and a breach-free run relaxes it — split conformal’s guarantee made robust to distribution drift, at the cost of targeting long-run rather than conditional coverage.

Table 5: Coverage of nominal-90% credible/confidence bands for the residual (VaR, ES), 150 CRSP names.

Method	VaR		ES	
	$\alpha = 1\%$	$\alpha = 2.5\%$	$\alpha = 1\%$	$\alpha = 2.5\%$
Block bootstrap (residuals)	0.947	0.953	0.86	0.90
i.i.d. bootstrap (residuals)	0.96	0.947	—	—
Gibbs, naive $\omega = 1$	0.987–0.993	(over-wide)	—	—
Gibbs, ω -calibrated	0.953–0.967		—	—
EVT/GPD ES interval	—	—	0.887	0.90

Reading guide — two different probability levels appear here, do not conflate them: α is the *risk level* of the quantile being estimated (VaR at 1% or 2.5%); each entry is the realized *coverage of the nominal-90% uncertainty band* around that risk number, i.e., the fraction of experiments in which the band contained the truth. Ideal is exactly 0.90: above 0.90 means the band is wastefully wide (over-conservative), below means overconfident. Bold marks the entry closest to 0.90 in each column — so the calibrated Gibbs (0.953–0.967) *beats* the naive one (0.987–0.993), whose larger numbers look reassuring but signal over-wide bands. The naive- ω Gibbs over-covers because pinball is shallow in the sparse tail; matching the posterior SD to the block SD repairs it. The 99% ES block interval under-covers (0.86, too few tail points); the GPD interval reaches 0.887 at the honest price of width (2.12 vs 1.31).

6.2 The repair: residual space

The identical construction on GARCH-standardized residuals gives $R = 0.79$: the overconfidence is gone (mildly conservative, the safe direction). The mechanism is confirmed independently: on residuals, block and i.i.d. bootstraps agree (VaR coverage 0.947 vs 0.953 at $\alpha = 2.5\%$) — the filter has absorbed the dependence, restoring the i.i.d.-like regime in which loss-based posteriors are honest. This is the same scale/shape division of labor as the residual-hybrid, now applied to *uncertainty* rather than accuracy. An adaptive conformal arm (Gibbs and Candès, 2021) corroborates from the coverage side: ACI attains the best and most stable realized VaR coverage under drift (5.03% realized at the 5% level; rolling max-deviation 0.042, vs 0.059 GARCH and 0.074 for a fixed empirical window). A complementary, benchmark-free reliability signal comes from ensemble disagreement: the dispersion of quantile forecasts across independently trained networks correlates +0.11 with the model’s own realized error — crude, but free, and usable as a first-line screen ahead of the formal bands of Section 6.3.

6.3 Calibrated credible bands on (VaR, ES)

Table 5 assembles the end-to-end result: block bootstrap on residuals for VaR and 2.5% ES; the ω -calibrated Gibbs posterior as the parametric credible band; EVT for the deep ES tail — every piece at or near nominal. To our knowledge no standard risk model (GARCH, CAViaR, FHS, or our own hybrid) supplies any counterpart. Two refinements remain registered: the loss-in-the-exponent ablation (CRPS, tail-CRPS, the Fissler–Ziegel joint (VaR, ES) score (Fissler and Ziegel, 2016), MEU) and formal comparison of ω -calibration to SafeBayes-style procedures (Grünwald and van Ommen, 2017).




Figure 6: Honest uncertainty, assembled. Left: the naive Gibbs posterior SD understates the block-bootstrap sampling SD $1.64\times$ on raw returns (i.i.d. bootstrap 1.19 — the gap is the dependence contribution); the ratio falls to 0.79 on GARCH residuals. Right: coverage of nominal-90% bands (bl = block bootstrap; V = VaR, E = ES at the stated level; Gb = Gibbs, naive ω vs ω -calibrated). Block bootstrap and the EVT interval sit at or near target; the naive Gibbs over-covers.

7 Likelihood-free inference for intractable volatility models

The final capability argument requires no comparison to GARCH because GARCH cannot enter: amortized posterior inference for models with intractable likelihoods, validated by simulation-based calibration.

Heston. From 40k simulated (parameter, 252-day path) pairs and ten path summaries, the amortized posterior is calibrated (central coverage cov_{80} 0.78–0.80, cov_{90} 0.85–0.90 for all five parameters) and informative where a single path identifies: vol-of-vol info-gain 36.5% (fraction of prior RMSE removed), long-run variance 13%, with mean-reversion and leverage weakly identified — the honest, expected pattern.

Rough Bergomi (the marquee result). The rough-volatility model (Bayer et al., 2016; Gatheral et al., 2018) — volatility whose path is driven by fractional noise controlled by the Hurst exponent H , with $H < 0.5$ producing paths jaggier than any standard model allows — has a non-Markovian fractional driver and no tractable likelihood; neither MLE nor any GARCH-family filter can fit it. The amortized posterior over (H, η, ρ, ξ_0) from a single 252-day path (30k sims) **recovers the Hurst roughness H with info-gain 0.42 and SBC cov_{90} 0.91** (IQN; GBM 0.89); η info-gain 0.30; all coverages 0.85–0.92. Recovering the parameter that *defines* rough volatility — notoriously difficult classically — with calibrated uncertainty from one path is, we believe, the strongest demonstration to date that GBC does something structurally new in finance. Here the neural IQN is not merely competitive but canonical: it is the estimator that *samples* the posterior; trees provide fixed quantiles only.




Figure 7: Likelihood-free calibration: information gain (fraction of prior RMSE removed by one 252-day path) per parameter, Heston (left, κ/ρ at midpoints of their reported weak ranges) and rough Bergomi (right). Honesty check by SBC coverage: Heston cov_{90} 0.85–0.90 across all five parameters; rough Bergomi 0.85–0.92, with the Hurst roughness H at 0.91 — informative *and* calibrated.




Figure 8: The deep-tail crossover on the most concentrated weighting schemes (top-1 and top-3 books; Gaussian/DCC top-3 drawn at the midpoint of its reported 0.018–0.019). Across *all* schemes Gaussian/DCC realizes 0.018–0.024 — roughly twice the nominal 0.01 — while the covariance-scale / nonparametric-shape hybrid is near or under target. At the moderate 5% tail the ordering reverses.

8 The multivariate tail: negatives, then a crossover

We report the negatives first, because they discipline the claim. Three direct nonparametric attacks on the portfolio tail all lose to multivariate GARCH: (i) a direct pooled portfolio-quantile model on a diversified 25-name basket (DCC/CCC — multivariate GARCH with dynamic (DCC) or constant (CCC) correlations across assets — benchmarks after Engle, 2002) (DCC 0.1709 < CCC 0.1719 < GBM 0.1752, ordering preserved in the crisis decile); (ii) a concentration sweep ($K = 2, \dots, 25$ by pairwise correlation) — the gap does not close as the basket concentrates; (iii) a neural weight-input IQN $Q(\tau \mid s, w)$ trained across the simplex — roughly $2\times$ worse than CCC/DCC at every weighting, because a network fed summary state never recovers the quadratic form $w'\Sigma w$ that GARCH models explicitly. The covariance *is* the object; we concede equity co-crash to explicit covariance models.

The rescue is, once again, the hybrid division of labor: DCC supplies the portfolio scale $\sigma_{w,t}$; a generative IQN learns only the shape of the standardized portfolio residual, conditioned on concentration (the Herfindahl index: the sum of squared portfolio weights, higher meaning more concentrated), state, and weights. Calibration is restored (PIT-KS 0.30 $\rightarrow$ 0.055–0.075 — a Kolmogorov–Smirnov test of whether probability-integral-transformed outcomes look uniform, i.e. whether the model is calibrated — comparable to Gaussian/DCC), and a clean crossover emerges: Gaussian/DCC is adequate at the moderate 5% tail (where the IQN is slightly conservative) but systematically *thin* in the deep 1% tail (cov₀₁ 0.018–0.024, roughly twice nominal breaches), while the hybrid is near nominal everywhere and best precisely on concentrated books (top-1 concentration: cov₀₁ 0.0085 vs target 0.01; Gaussian 0.019). The multivariate contribution is thus representational and calibrational: a samplable joint downside law whose deep co-crash tail is honest where the industry model’s is thin — reported alongside, not instead of, the three negatives.

9 Applications

Forward-looking, not just backward-looking — but a monitor, not a switch. Because the frontier score is a trailing window on GARCH-standardized residuals, it is available *ex ante*: a desk computes today’s score from today’s data and acts on it for tomorrow’s risk number. What the evidence does *not* support is day-level model-switching. As Section 4.3 shows, the per-day oracle gap is unreachable (predicted-edge correlation +0.16), because which day a tail event lands is not forecastable from prior-day state, and a learned gate collapses to always-nonparametric. The edge is, in effect, negligible over the calm $\sim 90\%$ of asset-days and concentrated (+2.7% pinball, DM 6.5) in the top misspecification decile and the most turbulent volatility regimes (an $\sim$ fourfold gap, Section 4.3). The deployable reading is therefore not “wait, then switch” but “run the robust engine always and read the score as a dial”: the nonparametric/residual-hybrid family *ties* the parametric benchmark on calm days and wins in stress, so a desk pays nothing to hold it and harvests the stress-period edge automatically. The one setting in which the score genuinely drives a switch is the per-name/standalone model, where nonparametric quantiles underperform on calm days; there the score gates, and its short detection latency (median zero, mean two to four days) costs little because the edge is back-loaded within a misspecification episode.

Regulatory capital. The hybrid+EVT+conformal engine is an internal-models VaR/ES candidate that tops the accuracy table while passing both exception tests at both levels; and because FHS/GARCH over-state ES_{97.5} (by 8–11% at one day, more at ten), a calibrated engine translates directly into capital efficiency without added breach risk.

A misspecification monitor. The frontier score is a live, per-asset dial: when a desk’s asset enters the top decile of residual-shape extremity, its standard VaR model is currently, locally wrong by $\sim 2\text{--}3\%$ of pinball with high confidence. This is deployable as model-risk surveillance exactly because it is causal and cheap (a trailing window on GARCH residuals).

Cold-start risk for new listings. The amortized model is the only entrant in this paper that can produce a conditional risk forecast for an asset with *no return history*: an IPO, a spin-off, a newly listed coin or ETF gets day-one quantiles from its characteristics and its first ticks, because the model was trained on the pooled (state, loss) pairs of hundreds of other names. The practical numbers: the amortized forecast beats own-history benchmarks at *every* listing age with the edge largest when young ($\sim 6\text{--}10\%$ of pinball in the first month of trading, full-scale panel), and below roughly 250 trading days there is no fitted per-asset alternative at all — a GARCH cannot be estimated. The ablation of Section 3.3 tells a desk exactly how the machine works as history accrues: characteristics carry the forecast at birth; the name’s own recent realized volatility takes over as it ages, because the model has learned one universal vol-state-to-quantile mapping rather than per-name parameters. Risk limits, margining, and market-making inventory for the newest names are the immediate use cases.

Where the edge trades — and on which side. The universes where the nonparametric tail wins — crisis FX, volatility indices, credit spreads, frontier equities — are the desks where tail-shape-aware sizing and hedging matter most. The *direction* of the trade matters as much as its location. The companion GRAFT-Q paper shows that *buying* model-flagged crash insurance loses (-1.7% per month, $t = -1.9$) even in the months the calibrated model correctly flags as elevated-tail: the crash premium survives conditioning on a correct physical-measure forecast, so a right-but-still-unprofitable buy signal is not a deployable trade. The deployable side is the other one — *warehousing* (selling) the premium — with the model supplying tail control rather than alpha on the mean: model-timed overlays improve tail outcomes, not average returns, consistent with the tail being the only place the model knows something a fixed-shape filter does not. The contribution to a premium-harvesting book is risk management (when to size down or step aside ahead of a genuine crash), not signal generation.

Scenario generation. The SBC machinery of Section 7 calibrates intractable scenario-generator models (rough vol, and by extension Hawkes and jump-diffusion) with posterior uncertainty on stressed paths — a principled input to CCAR/ICAAP-style stress design.

10 Conclusion and limitations

A fixed-shape parametric filter is hard to beat exactly where its assumption holds, which is most assets on most days. The paper’s answer to “when and why beat it” is a measurable frontier: local violation of the residual-shape assumption, at whose top decile the distribution-free family wins decisively, significantly, and portably across four universes. The engine that harvests this — parametric scale, nonparametric shape, EVT tail, conformal finish — meets and exceeds the FRTB standard banks actually run. And the generative members of the family contribute what no point model can: calibrated uncertainty on the risk number, and likelihood-free posteriors for models whose likelihoods do not exist in closed form.

Limitations, stated plainly. Gradient-boosted trees, not the neural network, are the strongest tabular estimator; the neural IQN earns its place through sampling and amortized posterior

inference, not point accuracy. The frontier score is necessary-not-sufficient (carbon/corn counterexamples; Baltic Dry wins at moderate kurtosis via limit-move dynamics). Credit and freight show large accuracy edges but clustered exceptions — unresolved. The stress analysis lacks a 2008 window pending WRDS 2000–2024 CRSP (registered). Intraday extension is blocked on tick-data access. The 99% ES credible band reaches 0.887 rather than 0.90, at $\sim 1.6\times$ width — the honest price of deep-tail uncertainty. And the per-day oracle gap of Section 4.3 should discipline the field’s enthusiasm for model-switching: most of the day-level advantage of any model pair is realization noise, not predictable structure.

Registered extensions within this paper’s scope, in priority order: the Fissler–Ziegel joint (VaR, ES) scoring of the FRTB battery and a certainty-equivalent (MEU) translation of the ES edge into basis points of capital; the loss-in-the-exponent ablation for the generalized-Bayes posterior; a leverage summary statistic to identify ρ in the SBC studies; Hawkes and jump-diffusion simulators as the next likelihood-free targets; vector quantile regression (Carlier–Chernozhukov–Galichon) as the principled multivariate transport map to succeed the scale/shape hybrid; the 2000–2024 WRDS panel for a 2008 stress window; and the intraday frontier once tick data is available. Job scripts for the first four are written and queued.

Data and code availability

A public replication package reproducing every table and figure in this paper — estimation code for all stages of the engine ((2)–(5)), the misspecification score (9), the FRTB battery, the generative-posterior and SBC routines, and the exact walk-forward, seeds, and hyperparameter grids referenced in the text — is archived at github.com/ANONYMIZED/gbc-downside with a citable snapshot at doi.org/10.5281/zenodo.XXXXXXX (both to be de-anonymized on acceptance). Return data derive from CRSP, Bloomberg, and OptionMetrics under the authors’ institutional licenses and cannot be redistributed; the package includes the code to rebuild every panel from those sources, the derived non-proprietary series, and a synthetic example dataset (`toy_example.py`) that runs the full pipeline end-to-end without licensed data. The pre-registered tuning grid is in `TUNING_GRID_PREREG.md`; deferred analyses flagged in the text (the Fissler–Ziegel joint (VaR, ES) re-scoring, the full GPC iteration, and the 2000–2024 stress window) are tracked as open issues in the repository.

References

- August A. Balkema and Laurens de Haan. Residual life time at great age. *Annals of Probability*, 2(5):792–804, 1974.
- Giovanni Barone-Adesi, Kostas Giannopoulos, and Les Vosper. VaR without correlations for portfolios of derivative securities. *Journal of Futures Markets*, 19(5):583–602, 1999.
- Basel Committee on Banking Supervision. Minimum capital requirements for market risk. Technical report, Bank for International Settlements, 2019.

- Christian Bayer, Peter Friz, and Jim Gatheral. Pricing under rough volatility. *Quantitative Finance*, 16(6):887–904, 2016.
- Pier Giovanni Bissiri, Chris C. Holmes, and Stephen G. Walker. A general framework for updating belief distributions. *Journal of the Royal Statistical Society: Series B*, 78(5):1103–1130, 2016.
- Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.
- Victor Chernozhukov, Iván Fernández-Val, and Alfred Galichon. Quantile and probability curves without crossing. *Econometrica*, 78(3):1093–1125, 2010.
- Peter F. Christoffersen. Evaluating interval forecasts. *International Economic Review*, 39(4):841–862, 1998.
- Will Dabney, Georg Ostrovski, David Silver, and Rémi Munos. Implicit quantile networks for distributional reinforcement learning. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, pages 1096–1105, 2018.
- Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.
- Paul Embrechts, Claudia Klüppelberg, and Thomas Mikosch. *Modelling Extremal Events for Insurance and Finance*. Springer, 1997.
- Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.
- Robert F. Engle. Dynamic conditional correlation: A simple class of multivariate generalized autoregressive conditional heteroskedasticity models. *Journal of Business & Economic Statistics*, 20(3):339–350, 2002.
- Robert F. Engle and Simone Manganelli. CAViaR: Conditional autoregressive value at risk by regression quantiles. *Journal of Business & Economic Statistics*, 22(4):367–381, 2004.
- Tobias Fissler and Johanna F. Ziegel. Higher order elicibility and Osband’s principle. *Annals of Statistics*, 44(4):1680–1707, 2016.
- Jim Gatheral, Thibault Jaisson, and Mathieu Rosenbaum. Volatility is rough. *Quantitative Finance*, 18(6):933–949, 2018.
- Isaac Gibbs and Emmanuel Candès. Adaptive conformal inference under distribution shift. In *Advances in Neural Information Processing Systems 34 (NeurIPS)*, pages 1660–1672, 2021.
- Lawrence R. Glosten, Ravi Jagannathan, and David E. Runkle. On the relation between the expected value and the volatility of the nominal excess return on stocks. *Journal of Finance*, 48(5):1779–1801, 1993.

- Tilman Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Peter Grünwald and Thijs van Ommen. Inconsistency of Bayesian inference for misspecified linear models, and a proposal for repairing it. *Bayesian Analysis*, 12(4):1069–1103, 2017.
- Peter R. Hansen, Asger Lunde, and James M. Nason. The model confidence set. *Econometrica*, 79(2):453–497, 2011.
- Wenxin Jiang and Martin A. Tanner. Gibbs posterior for variable selection in high-dimensional classification and data mining. *Annals of Statistics*, 36(5):2207–2231, 2008.
- Roger Koenker and Gilbert Bassett. Regression quantiles. *Econometrica*, 46(1):33–50, 1978.
- Hans R. Künsch. The jackknife and the bootstrap for general stationary observations. *Annals of Statistics*, 17(3):1217–1241, 1989.
- Paul H. Kupiec. Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives*, 3(2):73–84, 1995.
- Alexander J. McNeil and Rüdiger Frey. Estimation of tail-related risk measures for heteroscedastic financial time series: An extreme value approach. *Journal of Empirical Finance*, 7(3–4):271–300, 2000.
- Maria Nareklashvili, Nicholas G. Polson, and Vadim Sokolov. Generative quantile Bayesian prediction. *arXiv preprint arXiv:2510.21784*, 2025.
- James Pickands. Statistical inference using extreme order statistics. *Annals of Statistics*, 3(1):119–131, 1975.
- Nicholas G. Polson and Vadim Sokolov. Generative AI for Bayesian computation. *arXiv preprint arXiv:2305.14972*, 2023.
- Yaniv Romano, Evan Patterson, and Emmanuel J. Candès. Conformalized quantile regression. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, pages 3538–3548, 2019.
- Nicholas Syring and Ryan Martin. Calibrating general posterior credible regions. *Biometrika*, 106(2):479–486, 2019a.
- Nicholas Syring and Ryan Martin. Gibbs posterior inference on value-at-risk. *Scandinavian Actuarial Journal*, 2019(7):548–557, 2019b.
- Sean Talts, Michael Betancourt, Daniel Simpson, Aki Vehtari, and Andrew Gelman. Validating Bayesian inference algorithms with simulation-based calibration. *arXiv preprint arXiv:1804.06788*, 2018.
- Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.

Gilles O. Zumbach. The RiskMetrics 2006 methodology. Technical report, RiskMetrics Group, 2007.

A Glossary

Written so that a reader from any one field can parse every other field’s terms. Formulas referenced are in Section 3.

VaR / ES Value-at-Risk: the loss threshold exceeded with probability τ (a quantile with the sign flipped); Expected Shortfall: the average loss *given* VaR is exceeded (§3.1). FRTB capitalizes ES at $\tau = 2.5\%$.

Pinball (check) loss The asymmetric absolute loss (1) whose minimizer is the true quantile; the strictly proper score for quantile forecasts. Reported edges (“+2.7% of pinball”) are percentage reductions in this loss.

GARCH- t / GJR / EWMA Parametric filters for conditional variance (2); GJR adds a leverage term (down-moves raise vol more); EWMA is the fixed-decay special case (RiskMetrics); Student- t fixes the innovation shape.

HS / FHS Historical simulation: empirical quantiles of past returns. Filtered HS: empirical quantiles of GARCH-standardized residuals, rescaled by today’s $\hat{\sigma}_t$ — the regulatory workhorse; the special case of (5) with no state dependence.

CAViaR Conditional autoregressive VaR: models the quantile itself as an autoregression, estimated on the pinball loss; the one non-nested rival.

EVT / GPD / POT Extreme-value theory: exceedances over a high threshold converge to the generalized Pareto distribution; fitting it (peaks-over-threshold) gives the tail extrapolation (4).

Conformal prediction Distribution-free finite-sample correction sized on a held-out split (Proposition 1); guarantees coverage under exchangeability without any distributional assumption.

Quantile regression / IQN Estimating conditional quantiles by pinball ERM; the implicit quantile network (6) takes the level τ as an *input*, so one network carries the whole distribution and sampling is evaluation at random τ .

GBM (trees) Gradient-boosted regression trees on the pinball loss; the strongest point estimator on tabular features; cannot sample.

Amortization Training one model across many assets so inference for a new asset is evaluation, not estimation; the generative-Bayes reading is the inverse map $z \stackrel{D}{=} H(s, \tau)$ (§3.3).

Gibbs / generalized-Bayes posterior A posterior built from a loss instead of a likelihood, (10); ω tempers how fast the loss overwhelms the prior.

SBC Simulation-based calibration: simulate known truths, check the posterior’s ranks are uniform; the audit in Algorithm 3.

Walk-forward, purge, embargo Refit on past data only, on a schedule; delete observations adjacent to the train/test boundary so overlapping windows cannot leak information.

DM test Diebold–Mariano: a t -test on the loss difference between two forecasters, robust to serial correlation.

MCS Model Confidence Set: the subset of models not statistically dominated at a given level — “co-best” means surviving in that set.

Kupiec / Christoffersen Regulatory exception tests: correct *number* of VaR breaches (unconditional coverage) and breaches arriving *independently* rather than in clusters.

DCC Dynamic conditional correlation: the standard multivariate GARCH benchmark for portfolio tails.

VRP; $\mathbb{P}$ vs. $\mathbb{Q}$ Variance risk premium: option-implied variance exceeds realized on average. $\mathbb{P}$ is the physical measure (what happens); $\mathbb{Q}$ the risk-neutral measure (what option prices imply); their wedge at the put wing is the crash premium studied in the companion paper.

Misspecification frontier The trailing residual-shape score (9) and the empirical fact that the nonparametric edge concentrates in its top decile.

B Proofs

The four formal results of Section 3.4 are stated with proof sketches in the body; the complete arguments, with all regularity hypotheses made explicit, are collected here. Throughout, $\rho_\tau(u) = u(\tau - \mathbb{I}\{u < 0\})$ and $q_F = F^{-1}(\tau) = \inf\{x : F(x) \geq \tau\}$ is the lower generalized inverse.

Proof of Lemma 1 (pinball regret identity)

Assume $\mathbb{E}|Z| < \infty$, so

$$S(q) = (1 - \tau) \int_{(-\infty, q]} (q - z) dF(z) + \tau \int_{(q, \infty)} (z - q) dF(z)$$

is finite for every q , and S is convex (an expectation of the convex maps $q \mapsto \rho_\tau(z - q)$). For fixed z , $q \mapsto \rho_\tau(z - q)$ is Lipschitz with constant $\max(\tau, 1 - \tau)$, so the difference quotients $[\rho_\tau(Z - q') - \rho_\tau(Z - q)]/(q' - q)$ are bounded by the integrable constant $\max(\tau, 1 - \tau)$; dominated convergence justifies differentiating under the expectation wherever the one-sided limits exist. Off the atoms of F ,

$$\frac{d}{dq} \rho_\tau(z - q) = \mathbb{I}\{z < q\} - \tau, \quad \implies \quad S'(q) = \mathbb{E} \mathbb{I}\{Z < q\} - \tau = F(q^-) - \tau,$$

which equals $F(q) - \tau$ for all but countably many q . Since S is convex, hence absolutely continuous, the fundamental theorem of calculus gives

$$S(q) - S(q_F) = \int_{q_F}^q (F(u) - \tau) du,$$

the countable atom set being Lebesgue-null; *no continuity hypothesis on F is needed for this identity*. For nonnegativity, note $q_F = \inf\{x : F(x) \geq \tau\}$ implies $F(u) \geq \tau$ for $u \geq q_F$ and $F(u) < \tau$ for $u < q_F$. If $q > q_F$ the integrand is ≥ 0 on $[q_F, q]$, so the integral is ≥ 0 ; if $q < q_F$,

$$S(q) - S(q_F) = \int_{q_F}^q (F(u) - \tau) du = \int_q^{q_F} (\tau - F(u)) du \geq 0,$$

the integrand again ≥ 0 on $[q, q_F]$. Equality holds iff $F(u) = \tau$ for Lebesgue-a.e. u between q_F and q . Finally, if F is absolutely continuous with density $0 < m \leq f \leq M$ on the interval joining q_F and q , then F is continuous there, so $F(q_F) = \tau$ and $F(u) - \tau = \int_{q_F}^u f$; hence $m|u - q_F| \leq |F(u) - \tau| \leq M|u - q_F|$, and integrating over the interval yields $\frac{m}{2}(q - q_F)^2 \leq S(q) - S(q_F) \leq \frac{M}{2}(q - q_F)^2$. $\square$

Proof of Proposition 1 (split-conformal validity)

Let $k = \lceil (n+1)\tau \rceil$ and assume $1 \leq k \leq n$ so that the k -th order statistic $e_{(k)} = c_\tau$ of the n calibration errors $\{e_1, \dots, e_n\}$ exists; when $k = n+1$ (i.e. τ within $1/(n+1)$ of 1) set $c_\tau = +\infty$, giving coverage 1. Suppose the $n+1$ errors $e_1, \dots, e_{n+1}$ (calibration plus the test error $e_{n+1} = z_{n+1} - \tilde{q}(\tau | s_{n+1})$) are exchangeable and almost surely distinct. Then the rank R of e_{n+1} among the pooled $n+1$ values is uniformly distributed on $\{1, \dots, n+1\}$. By construction $\tilde{q}(\tau | s_{n+1}) = \hat{q}(\tau | s_{n+1}) + c_\tau$, so

$$\{z_{n+1} \leq \tilde{q}(\tau | s_{n+1})\} = \{e_{n+1} \leq c_\tau\} = \{e_{n+1} \leq e_{(k)}\} = \{R \leq k\},$$

the middle equality using a.s. distinctness (with ties the middle equality fails; e.g. if all e_i are equal the event has probability 1, which can exceed the stated upper bound). Therefore

$$\Pr(z_{n+1} \leq \tilde{q}(\tau | s_{n+1})) = \Pr(R \leq k) = \frac{k}{n+1} = \frac{\lceil (n+1)\tau \rceil}{n+1}.$$

Since $(n+1)\tau \leq \lceil (n+1)\tau \rceil < (n+1)\tau + 1$, dividing by $n+1$ gives $\tau \leq k/(n+1) < \tau + \frac{1}{n+1}$. The probability is over the joint law of calibration and test points; the guarantee is marginal, not conditional on s_{n+1} , because a single scalar shift c_τ pooled across states cannot deliver state-conditional coverage. $\square$

Proof of Proposition 2 (tail wedge)

Fix $\nu > 2$, so the t_ν law has finite variance and its unit-variance version $Z = T_\nu \sqrt{(\nu-2)/\nu}$ is well defined. The t_ν density is $\sim C_\nu |x|^{-(\nu+1)}$ as $|x| \rightarrow \infty$, so the lower tail is regularly

varying with index $-\nu$: $\Pr(T_\nu < -x) = \tilde{c}_\nu x^{-\nu}(1 + o(1))$. Unit-variance scaling multiplies the constant, $\Pr(Z < -x) = c_\nu x^{-\nu}(1 + o(1))$ with $c_\nu = \tilde{c}_\nu((\nu - 2)/\nu)^{\nu/2}$, and leaves the index ν unchanged. Writing $U(\tau) = |Q_\nu(\tau)|$ and inverting the regularly varying tail (Embrechts et al., 1997) gives $\tau = c_\nu U(\tau)^{-\nu}(1 + o(1))$, hence $U(\tau) = (c_\nu/\tau)^{1/\nu}(1 + o(1))$ as $\tau \downarrow 0$, i.e. $Q_\nu(\tau) = -(c_\nu/\tau)^{1/\nu}(1 + o(1))$. For $\nu_1 < \nu_2$,

$$\frac{|Q_{\nu_1}(\tau)|}{|Q_{\nu_2}(\tau)|} = \frac{(c_{\nu_1}/\tau)^{1/\nu_1}}{(c_{\nu_2}/\tau)^{1/\nu_2}} (1 + o(1)) = \text{const} \cdot \tau^{-(1/\nu_1 - 1/\nu_2)} (1 + o(1)) \rightarrow \infty$$

because $1/\nu_1 > 1/\nu_2$. Thus for all sufficiently small τ both quantiles are negative with $|Q_{\nu_1}(\tau)| > |Q_{\nu_2}(\tau)|$, i.e. $Q_{\nu_1}(\tau) < Q_{\nu_2}(\tau)$, and the wedge $|Q_{\nu_1}(\tau) - Q_{\nu_2}(\tau)| \rightarrow \infty$. The ordering therefore holds for every $\tau < \delta$ for some $\delta > 0$; putting $\tau^* = \min(\delta, \frac{1}{2})$ gives $\tau^* \in (0, \frac{1}{2})$ as asserted. (This establishes *existence* of τ^* ; the actual crossover τ_c where the ordering flips is not claimed here and is located numerically in Remark 1.) Under local truth t_{ν_1} , applying (8) with $q = Q_{\nu_2}(\tau)$ and $q_F = Q_{\nu_1}(\tau)$ bounds the regret below by $\frac{m}{2}(Q_{\nu_1}(\tau) - Q_{\nu_2}(\tau))^2$, with m the minimum of the t_{ν_1} density on the wedge interval. $\square$

Proof of Lemma 2 (generative sampler)

Let $\tilde{q}(\cdot | s) : (0, 1) \rightarrow \mathbb{R}$ be nondecreasing and left-continuous, $\tau \sim U(0, 1)$, and $Z^* = \tilde{q}(\tau | s)$. Fix $x \in \mathbb{R}$ and put $G(x) = \sup\{t \in (0, 1) : \tilde{q}(t | s) \leq x\}$ (with $\sup \emptyset = 0$). Because $\tilde{q}$ is nondecreasing, $\{t : \tilde{q}(t | s) \leq x\}$ is an initial subinterval of $(0, 1)$; left-continuity of $\tilde{q}$ makes the supremum attained, so the set is $(0, G(x)]$ and

$$\{\tilde{q}(\tau | s) \leq x\} = \{\tau \leq G(x)\}.$$

Since $\tau \sim U(0, 1)$, $\Pr(Z^* \leq x) = \Pr(\tau \leq G(x)) = G(x)$; thus Z^* has cumulative distribution function G , the generalized inverse of $\tilde{q}$. For a nondecreasing left-continuous $\tilde{q}$ the generalized inverse of its generalized inverse recovers the original function, so the quantile function of Z^* is exactly $\tilde{q}(\cdot | s)$. $\square$